Saying less, steering more

Simulation 2 (asynchronous activation), setup `task003-nosolution`: first results

LLM-free simulator; extended with the efficiency of control, 12 October 2026; runs of 9 October 2026

How the simulator works. The world, `task_003`, has three candidate allocations; A0 is the correct one. Each of 24 agents is given one true fact about it and reasons as an exact Bayesian: its *belief* is the exact probability of A0, A1 and A2 given the facts it currently remembers. There are no rounds. Each agent acts at random times, on average once per time unit: it reads q posts drawn at random from the 48 most recent posts on a shared *board* (never its own), adds their facts to its memory, recomputes its belief, *votes* by probability matching (a vote drawn from its belief) and posts one remembered fact that supports its vote, or nothing if none does. Every remembered fact is forgotten after a random lifetime (on average a fraction ρ survives one time unit; reading it again renews it). A *controller* with a target, A0 (truth) or A2 (false), acts at random times at rate λ_c until $t = 30$: it reads q_c recent agent posts and, unless the *vote gate* (at least 75% of them from agents voting for its target) or the *proof stop* (its target already proved) silences it, posts the one fact from its pool that best supports its target, at most 30 messages per episode. Agents continue alone until $t = 40$. Every episode is run three times with the same random draws: without a controller, with the truth controller and with the false controller.

Summary. 24 agents hold evidence that, pooled, is a coin flip between allocation A0 (the truth) and A2, with no proof available to anyone. Simulation 2 removes the day/night rounds of Simulation 1: agents and the controller act at random times, remembered facts fade continuously, and the shared board fades too. Both controllers still tip the swarm, and **the effect persists** after they stop. When agents forget, control is **weaker than in Simulation 1** (truth 0.29 against 0.42; false 0.19 against 0.28); with full memory it is the same. The controller's rate barely matters when agents forget. With full memory, **a faster controller steers less**: it puts more of its pool into circulation, and the whole pool pulls every belief toward the same point whatever the target. Turning the vote gate off adds about 16 messages per episode and no control. Main hypotheses were confirmed on a second, independent set of random draws.

1 Setup

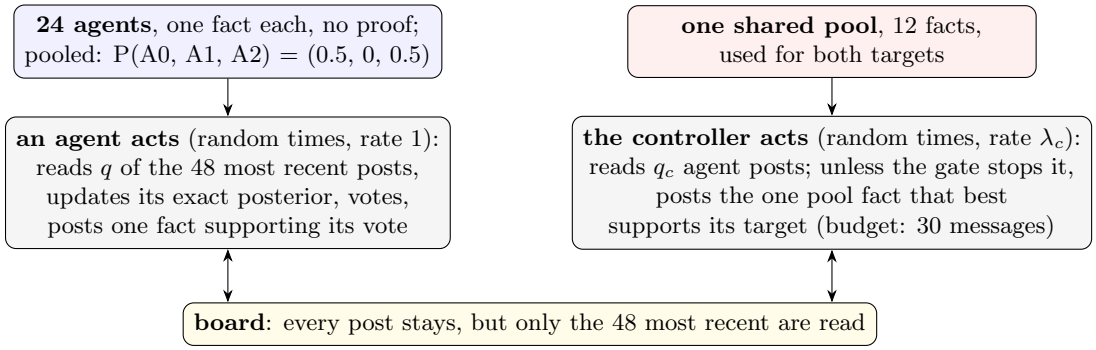


Diagram 1. Simulation 2. There are no rounds. Each remembered fact is forgotten after a random lifetime (on average a fraction ρ survives one time unit) and comes back if read again. The controller acts for $0 \le t < 30$; agents continue alone until $t = 40$.

The agents' facts plus the *whole* pool give P(A0, A1, A2) = (0.75, 0, 0.25): the truth keeps an edge that cannot be designed away, because every fact is true. Agents are exact Bayesian reasoners and vote by probability matching (the vote is drawn from the posterior). The *vote gate* keeps the controller silent when at least 75% of the posts it read come from agents who voted for its target. The grid: posts read $q \in \{3, 6, 12\}$, $q_c \in \{3, 6, 12, 24\}$, memory $\rho \in \{0.75, 1\}$, controller rate $\lambda_c \in \{0.5, 1, 2, 4, 8\}$ actions per time unit, gate on/off; 1,000 episodes per setting. Control is measured by the *gain*

$$G(t) = \frac{1}{E}\sum_{e=1}^{E}\Big[m_{\text{target}}(t)^{\text{ctrl}}(e) - m_{\text{target}}(t)^{\text{silent}}(e)\Big], \tag{1}$$

where $m_k(t)$ is the agents' mean belief in allocation k and each episode e uses the same random draws with and without the controller. Unless stated otherwise: $q = 6$, $q_c = 12$, gate on.

2 Main result: both controllers tip the swarm, and the effect stays

Table 1: Gain at $t = 30$, Eq. (1), $\lambda_c = 1$, gate on; in brackets the silent swarm's belief in the target. Simulation 1: one fact per night, same q, q_c, ρ.

ρ	controller	Simulation 2	Simulation 1	messages sent
0.75	truth (A0)	0.29 (0.38)	0.42	11
0.75	false (A2)	0.19 (0.48)	0.28	12
1	truth (A0)	0.38 (0.50)	0.38	7
1	false (A2)	0.14 (0.50)	0.15	15

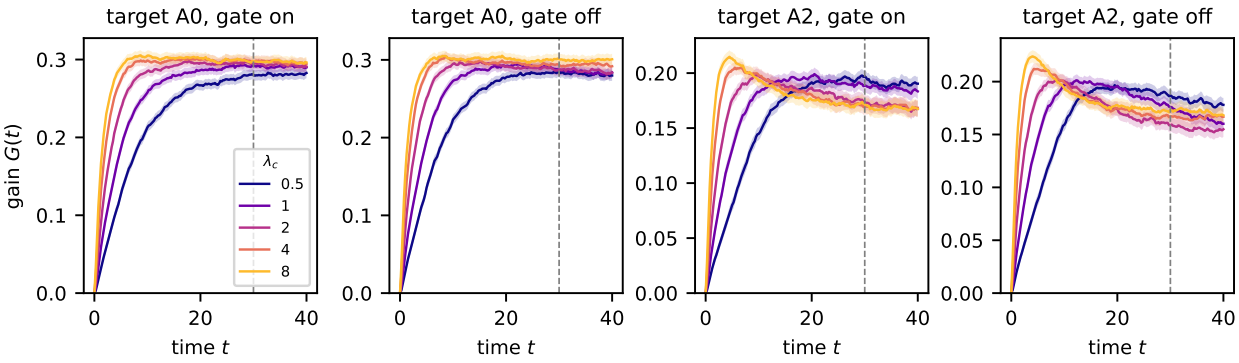


Figure 1: Gain over time by controller rate λ_c, $\rho = 0.75$; the controller stops at $t = 30$ (dashed). Bands: 95% intervals.

- **Both controllers tip the swarm** (Table 1); the truth controller is stronger, as in Simulation 1.
- **When agents forget, control is weaker than in Simulation 1**, by 0.14 (truth) and 0.09 (false) with fresh random draws. With full memory the two simulations agree within 0.01. In Simulation 1 memories change only at dawn and every agent acts right after a full day of posts; here facts expire at random moments between actions, so fewer of the controller's facts are held together at any time.
- **The effect persists**: ten time units after the controller stops, the gain has changed by at most 0.016 at $\rho = 0.75$ (Figure 1); agents keep relaying the controller's facts (12–45% of their posts).

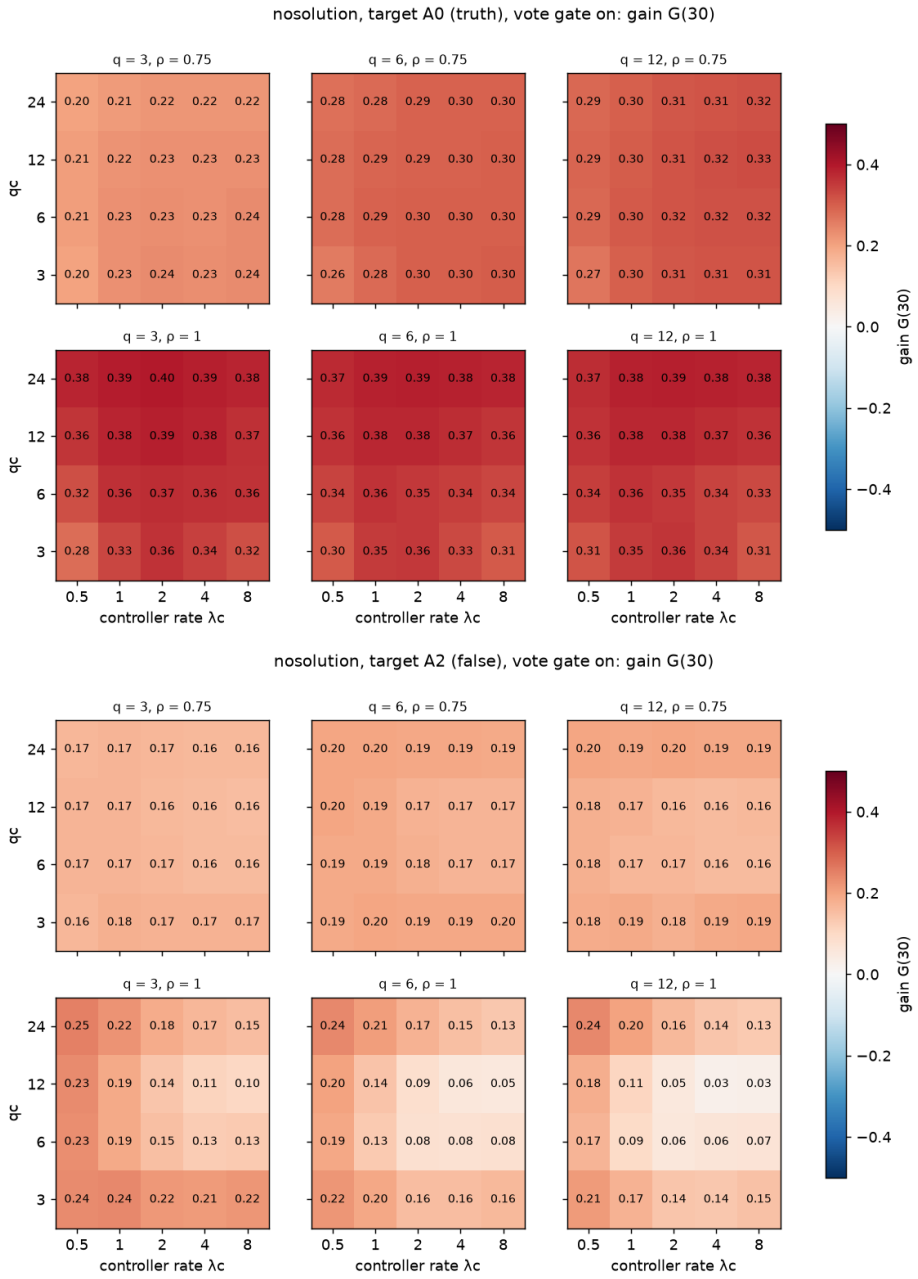


Figure 2: Gain $G(30)$ over controller rate (x) and posts read by the controller q_c (y), for every q (columns) and ρ (rows); gate on. Top: target A0; bottom: target A2.

3 The controller's rate

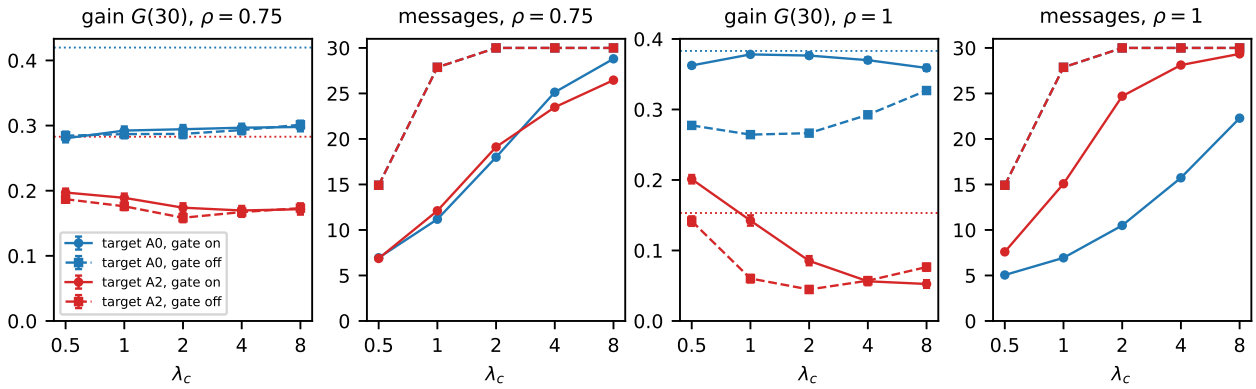


Figure 3: Gain at $t = 30$ and messages sent against the controller rate, for $\rho = 0.75$ and $\rho = 1$ (target A0 blue, A2 red; gate on solid, off dashed; dotted: Simulation 1).

When agents forget, the rate barely matters. From $\lambda_c = 0.5$ to 8 the gain moves from 0.28 to 0.30 (truth) and from 0.20 to 0.17 (false), but the messages sent rise from about 7 to 26–29. A fast controller spends its budget early (by $t \approx 4$ with the gate off), reaches its gain sooner, and then the gain stays put. Gain per message is highest for a slow, gated controller.

With full memory, a faster controller steers less. The false controller's gain falls from 0.20 at $\lambda_c = 0.5$ to 0.05 at $\lambda_c = 8$, and with the gate off the truth controller loses 0.05 on average against the gate on, up to 0.12 (Figure 3, right). The same fall appears at every q and at $q_c = 3$, 6 and 12 (Figure 2, bottom rows).

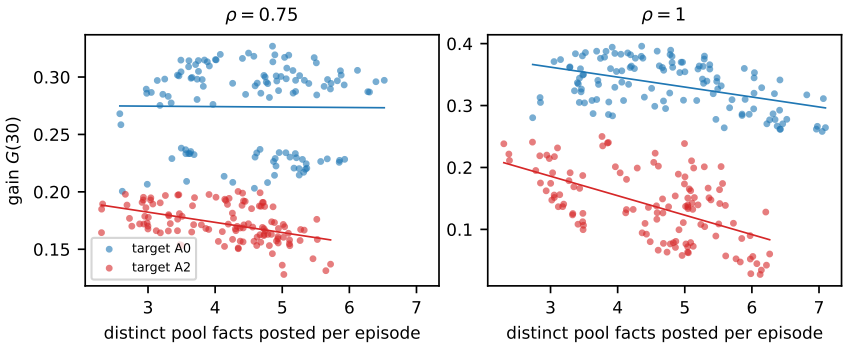


Figure 4: Each point is one setting (q, q_c, λ_c, gate): gain at $t = 30$ against the number of different pool facts the controller posted per episode. Lines: least-squares fits.

Why: saying more of the pool pulls toward the same point. Every fact is true, and the agents' facts plus the whole pool give (0.75, 0, 0.25). With full memory nothing posted is forgotten, so the more *different* pool facts a controller puts into circulation, the closer beliefs move to that point, whatever its target: down from the false controller's side, and down from the truth controller's best, which lies above 0.75. At $\rho = 1$ each additional distinct fact lowers the gain by 0.031 (false) and 0.016 (truth) (Figure 4, right). The rate sets that number: spread over 30 time units, the controller reads varied boards and posts more different facts (6.3 per episode at $\lambda_c = 1$, gate off); at $\lambda_c = 8$ it spends its budget in a burst on similar boards and repeats itself (5.0), which is why the gain partly recovers there. When agents forget, facts leave again and the effect is weak (Figure 4, left). This was found after looking at the data, so it was rerun with fresh random draws: **every one of the 16 checked reversals reappeared** (false controller, gain at $\lambda_c = 0.5 \ldots 8$: 0.20, 0.15, 0.08, 0.05, 0.05). The explanation by distinct facts rests on a correlation across settings and is not yet tested directly (for example with a controller that remembers what it has said).

4 The vote gate

The gate keeps the controller silent when it reads enough support. Turning it off adds about 16 messages per episode at $\lambda_c = 1$ (confirmed with fresh draws) and **no control** when agents forget (-0.017 and -0.003); with full memory it *lowers* the gain (truth -0.05 on average) through the mechanism above.

5 The controller overestimates its support

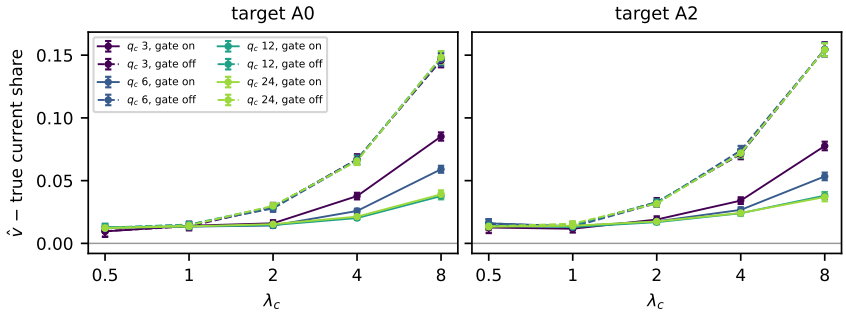


Figure 5: Error of the controller's estimate $\hat{v}$ (share of posts read whose author voted for the target) against the true current share of all 24 agents' last votes; $\rho = 0.75$; solid gate on, dashed off.

Unlike in Simulation 1, $\hat{v}$ is biased upward, and **the faster the controller, the larger the bias**: about $+0.014$ at $\lambda_c \leq 1$, rising to $+0.05$ (gate on) and $+0.15$ (gate off) at $\lambda_c = 8$. A likely reason, not yet tested: the posts it reads come from agents who acted recently, often just after reading the controller's facts, while the true share also counts older votes. More posts read (q_c) reduce the error's spread (0.16 at $q_c = 3$, 0.07 at $q_c = 24$) but not the bias.

6 How efficiently does the controller steer?

A controller can change the swarm in two ways: in the direction it is asked to, or in ways that happen whatever its target (for example pulling agents out of A1, or splitting them). We call the share of the first kind its *efficiency*. The target is chosen by a fair coin, so how much the swarm's state reveals about it is a well-defined information.

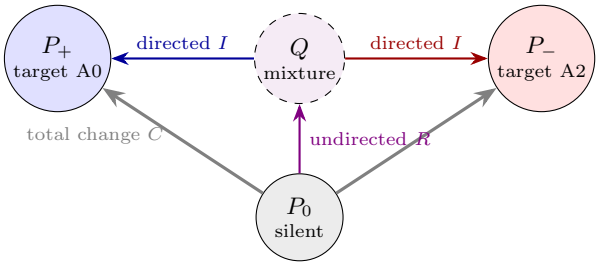


Diagram 2. Three distributions of the swarm's state at time t, one per arm, each over the episodes of that arm. The controller's total change C splits exactly into a part that depends on its target (I, from the mixture Q to each arm) and a part that does not (R, from silence to Q). Efficiency $\eta = I/C$ is the directed share.

For each arm a (silent 0, target A0 $+$, target A2 $-$) and a fixed time t, each episode e contributes a probability vector $y_a(e)$ over the values x of the state X, and

$$\hat{P}_a(x) = \frac{1}{n}\sum_{e=1}^{n} y_a(e)(x), \qquad I = \tfrac{1}{2}\,\mathrm{KL}(\hat{P}_+\|\hat{Q}) + \tfrac{1}{2}\,\mathrm{KL}(\hat{P}_-\|\hat{Q}), \qquad C = \tfrac{1}{2}\,\mathrm{KL}(\hat{P}_+\|\hat{P}_0) + \tfrac{1}{2}\,\mathrm{KL}(\hat{P}_-\|\hat{P}_0), \quad (2)$$

with $\hat{Q} = \frac{1}{2}(\hat{P}_+ + \hat{P}_-)$ and $\mathrm{KL}(p\|q) = \sum_x p(x)\log_2[p(x)/q(x)]$. Exactly $C = I + \mathrm{KL}(\hat{Q}\|\hat{P}_0)$, and $\eta = I/C \in [0,1]$. Two resolutions: **one agent** (X = the vote of a randomly chosen agent; $y_a(e)$ = that episode's vote shares) and the **swarm** ($X = s = m_{\mathrm{A0}} - m_{\mathrm{A2}}$, the difference of the agents' mean beliefs, in three equal bins on $[-1,1]$; $y_a(e)$ marks the episode's bin). At swarm level the controlled swarm often reaches states the silent swarm never visits, which makes C infinite; there we use the bounded version $\eta_{JS} = I/(I + \mathrm{JS}(\hat{Q},\hat{P}_0))$, with JS the Jensen–Shannon divergence. We also measure how unsure the agents become: $\bar{H}$, the mean entropy of the agents' beliefs (0 to 1.58 bits), and the disagreement $\mathrm{JSD} = H(\text{mean belief}) - \bar{H}$, each as controlled minus silent in the same episode ($\Delta\bar{H}$, $\Delta\mathrm{JSD}$). All quantities are taken at one time, across episodes; they are not divergences between whole trajectories (those would be at least as large). Setting: probability matching, $q = 6$, $q_c = 12$; 1,000 episodes per arm.

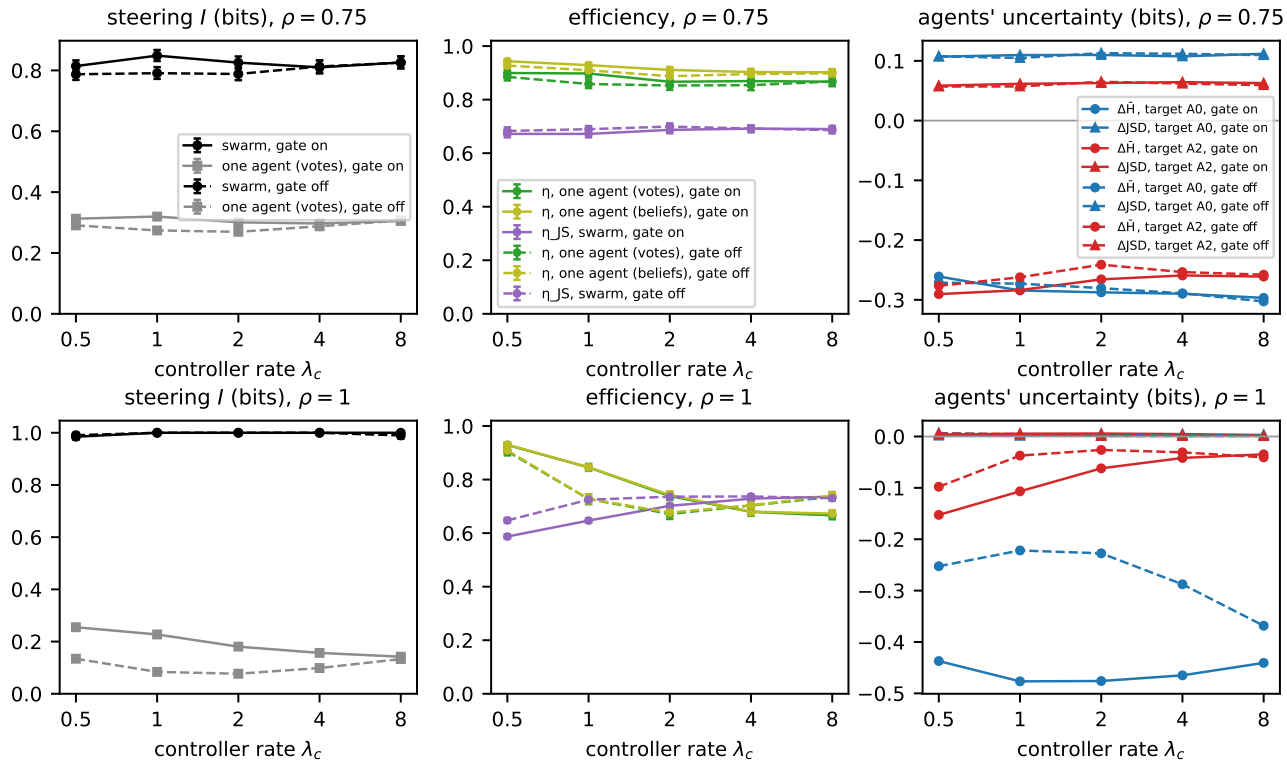


Figure 6: Steering I (left), efficiency (middle) and the change in the agents' uncertainty (right; circles $\Delta\bar{H}$, triangles ΔJSD) at $t = 30$ against the controller rate λ_c (solid: gate on; dashed: gate off). Bars: 95% intervals.

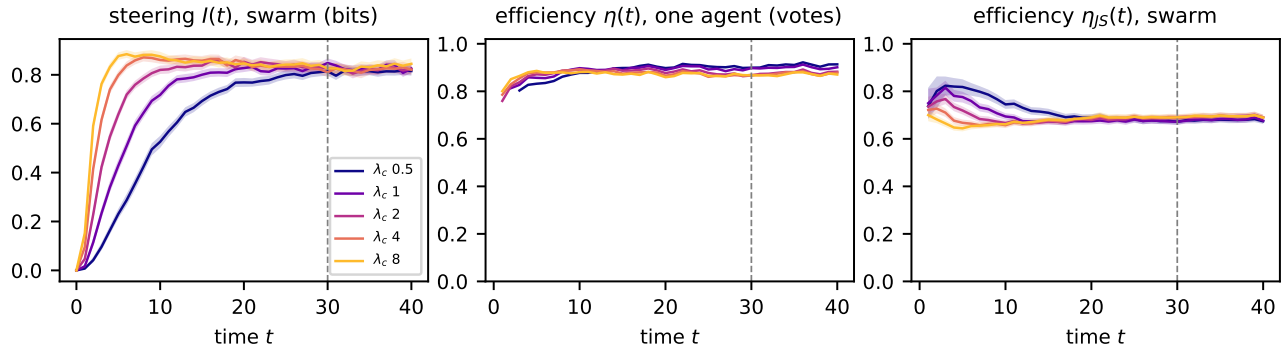


Figure 7: Over time, $\rho = 0.75$, gate on, one line per rate; the controller stops at $t = 30$ (dashed). Bands: 95% intervals.

Findings. (Figures 6, 7.)

- **When agents forget, control is efficient at every rate.** The swarm reveals the target well ($I \approx 0.8$ bits; one agent ≈ 0.3) and about 90% of the change is directed ($\eta = 0.85$–0.90), almost independent of the rate and of the gate. Fast controllers reach this sooner (Figure 7).
- **With full memory, faster control is less efficient**: one-agent η falls from 0.93 ($\lambda_c = 0.5$, gate on) to 0.67 ($\lambda_c = 8$), and turning the gate off lowers it further at slow rates. This is the mechanism of Section 3 (more different pool facts in circulation pull both targets toward the same point), now measured as undirected change.
- **Controllers make agents surer** ($\Delta\bar{H} \approx -0.27$ bits at $\rho = 0.75$) but split them somewhat (ΔJSD +0.06 for target A2, +0.11 for target A0).
- **The swarm-level I saturates at 1 bit** with full memory: every episode ends on its controller's side. I measures distinguishability, the gains of Section 3 measure size.
- **100 paths**: errors of ± 0.02–0.03 for I and η, ± 0.006 for ΔJSD (Table 2).

Table 2: Precision. Value from 1,000 paths per arm, and root-mean-square error of the estimate from 50, 100 or 300 paths (200 random subsets each), at $t = 30$, $\lambda_c = 1$, gate on, $\rho = 0.75$. *: share of subsets in which the quantity is infinite or undefined.

quantity	1,000 paths	error, 50	error, 100	error, 300
steering I, swarm (3 bins)	0.849	0.039	0.026	0.013
steering I, one agent (votes)	0.320	0.023	0.017	0.009
efficiency η, one agent (votes)	0.898	0.032	0.022	0.010
efficiency η_{JS}, swarm	0.672	0.027	0.019	0.010
$\Delta\bar{H}$, target A0	-0.284	0.024	0.016	0.009
$\Delta\bar{H}$, target A2	-0.284	0.023	0.016	0.008
ΔJSD, target A0	0.110	0.008	0.006	0.003
ΔJSD, target A2	0.062	0.008	0.006	0.003

Caveats. One world (task_003), 24 agents, exact-Bayesian agents, probability-matching votes; a controller without memory that posts one fact at a time; a budget of 30 messages. Means over 1,000 episodes per setting. Full analysis, code and tables: `analysis/task003_llm_free/simulation_2/analysis/` (report of 9 October, reviewer guide).