

ON THE CONTROL EFFICIENCY OF MULTI-AGENT LANGUAGE MODEL SYSTEMS

Anonymous authors

Paper under double-blind review

ABSTRACT

How efficiently can an external controller steer a group of interacting language-model agents? We study this question in relational reasoning tasks where evidence is distributed across the population. Our feedback controller has two separate limits: it observes only a sample of the agents and can intervene in only a fixed number of their interactions. This setup lets us measure the distinct effects of sensing and actuation. We introduce finite-time measures for three aspects of control: how much an intervention changes the population, how efficiently action-dependent information becomes coherent population motion, and how efficiently the finite-time thermodynamic control budget becomes directed current. We derive a finite-state reference theory with a single measurable directional affinity h and a kinetic compliance γ , which separates which direction controlled revisions favor from how readily agents revise at all. In experiments with matched initial social and knowledge states, advocacy increases support for the controller’s target in every tested condition, with susceptibilities from 0.100 to 0.163. The transmitted control information depends strongly on the population state. More actuation produces a wider informative region, whereas stronger feedback gain is predicted to concentrate information near the controller’s decision boundary. Information–response efficiency increases with actuation across the full resource grid, from about 0.4% to 2.5% of the Pinsker upper bound. A calibrated finite-compliance model also reproduces the observed controlled current closely ($r = 0.999$). Finally, the single-affinity theory yields the exact finite-time identity $\Delta S_{\text{sys}} + hJ_c + I(n_k; Y_k) = \Sigma \geq 0$ and the bounded thermodynamic efficiency $\eta_{\text{th}} = hJ_c/[hJ_c + I(n_k; Y_k)]$. The resulting theory predicts diminishing returns from additional sensing whenever the acquired information does not produce a proportional increase in directed current. Our framework measures control efficiency while keeping it separate from whether the agents reach the correct answer.

1 INTRODUCTION

OpenAI’s resolution of the long-standing Navier–Stokes Millennium Problem and the recent Hugging Face hacking incident are emblematic events of AI’s turn toward multi-agent systems and swarms (OpenAI, 2026; Greenblatt et al., 2026). From scientific discovery to agent-based social networks such as Moltbook, interacting populations of agents are becoming a distinct locus of capability in AI (Jiang et al., 2026). Crucially, these capabilities need not be reducible to those of the agents in isolation: interaction can give rise to collective behavior, conventions, biases, and higher-order coordination that are absent at the level of individual agents (Ashery et al., 2025; Riedl, 2026; Anthropic, 2026). This has prompted recent calls for a *science of agents collective behavior* concerned with collective capabilities, population-level failures, oversight, and control (Google DeepMind et al., 2026). In this view, intelligence is increasingly not only a property of the agent, but also of the interaction through which agents communicate and coordinate (Chopra, 2026).

Collective behavior has long been studied through statistical physics, game theory, and complex-systems approaches to emergence (Castellano et al., 2009). These traditions, however, have typically represented agents through compressed states, actions, payoffs, or local update rules, leaving the epistemic content of interaction—and the reasoning induced by that content—largely outside the model. Recent work has begun precisely the translation required for language-model populations: from reformulating rational cooperation and equilibrium through similarity inference in foundation-

model agents (Meulemans et al., 2026), to extending statistical-mechanical descriptions to collective regimes (El et al., 2026) and information-theoretic descriptions to higher-order coordination (Riedl, 2026). These developments broaden the classical study of collective behavior, but leave a different question largely open: once such collective dynamics arise, how can they be steered?

Much of this literature asks when coordination, consensus, or collective regimes emerge; we instead take control as the starting point. The Hugging Face incident makes this shift difficult to ignore: agents used a shared, unsanctioned message board to coordinate workstreams and redirect the activity of other agents, turning the communication substrate itself into a channel through which collective behavior could be influenced (Greenblatt et al., 2026). Whenever we study multi-agent populations or swarms, we are inevitably dealing with stochastic systems. The central scientific question is therefore not simply whether control succeeds, but how directed intervention competes with the endogenous fluctuations of the population. It is this competition that defines the regimes of stochastic feedback control we study and motivates our notion of *control efficiency*.

To study this problem, we construct a language-mediated population in which agents reason over distributed evidence and interact through a shared message board. Each agent carries both a social state—its current decision—and an epistemic state—the evidence available for subsequent reasoning. Into this population we introduce a resource-constrained feedback controller. Before each round, the controller observes only a finite sample of agents and, conditional on this partial observation, may intervene in only a bounded number of interactions. At the coarse population level, the feedback loop is

$$n_k \longrightarrow Y_k \longrightarrow U_k \longrightarrow n_{k+1}, \quad (1)$$

where n_k denotes current support for the controller target, Y_k its partial observation, and U_k its intervention. Finite sensing and finite actuation therefore constitute separate control resources. In our main setting, the controller communicates a recommendation but no task evidence: it can perturb the social dynamics without directly increasing the epistemic content of the population. This allows us to distinguish changes in what agents believe from changes in what they know, and to separate controllability from task correctness.

We characterize this feedback process along three complementary axes. *Susceptibility* measures the directed change in population state induced by intervention, separating systematic response from the population’s background dynamics. *Action-to-population information* measures how strongly the controller’s action is expressed in the distribution of subsequent population states; relative to the conditional entropy of the controller action, it quantifies how much of the available control information reaches the population. Finally, an *information–response efficiency*, obtained through a Pinsker-type bound, compares the information required to sustain the observed directed response with the action-dependent information actually transmitted to the population. Together, these quantities distinguish whether control is strong, whether it is statistically expressed above endogenous fluctuations, and whether the transmitted control information is efficiently converted into coherent population motion. They provide a finite-time account of how limited sensing and bounded intervention steer a reasoning population whose epistemic state continues to evolve.

2 RELATED WORK

Fusce mauris. Vestibulum luctus nibh at lectus. Sed bibendum, nulla a faucibus semper, leo velit ultricies tellus, ac venenatis arcu wisi vel nisl. Vestibulum diam. Aliquam pellentesque, augue quis sagittis posuere, turpis lacus congue quam, in hendrerit risus eros eget felis. Maecenas eget erat in sapien mattis porttitor. Vestibulum porttitor. Nulla facilisi. Sed a turpis eu lacus commodo facilisis. Morbi fringilla, wisi in dignissim interdum, justo lectus sagittis dui, et vehicula libero dui cursus dui. Mauris tempor ligula sed lacus. Duis cursus enim ut augue. Cras ac magna. Cras nulla. Nulla egestas. Curabitur a leo. Quisque egestas wisi eget nunc. Nam feugiat lacus vel est. Curabitur consectetur.

Suspendisse vel felis. Ut lorem lorem, interdum eu, tincidunt sit amet, laoreet vitae, arcu. Aenean faucibus pede eu ante. Praesent enim elit, rutrum at, molestie non, nonummy vel, nisl. Ut lectus eros, malesuada sit amet, fermentum eu, sodales cursus, magna. Donec eu purus. Quisque vehicula, urna sed ultricies auctor, pede lorem egestas dui, et convallis elit erat sed nulla. Donec luctus. Curabitur

et nunc. Aliquam dolor odio, commodo pretium, ultricies non, pharetra in, velit. Integer arcu est, nonummy in, fermentum faucibus, egestas vel, odio.

Sed commodo posuere pede. Mauris ut est. Ut quis purus. Sed ac odio. Sed vehicula hendrerit sem. Duis non odio. Morbi ut dui. Sed accumsan risus eget odio. In hac habitasse platea dictumst. Pellentesque non elit. Fusce sed justo eu urna porta tincidunt. Mauris felis odio, sollicitudin sed, volutpat a, ornare ac, erat. Morbi quis dolor. Donec pellentesque, erat ac sagittis semper, nunc dui lobortis purus, quis congue purus metus ultricies tellus. Proin et quam. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Praesent sapien turpis, fermentum vel, eleifend faucibus, vehicula eu, lacus.

3 BACKGROUND

Pellentesque habitant morbi tristique senectus et netus et malesuada fames ac turpis egestas. Donec odio elit, dictum in, hendrerit sit amet, egestas sed, leo. Praesent feugiat sapien aliquet odio. Integer vitae justo. Aliquam vestibulum fringilla lorem. Sed neque lectus, consectetur at, consectetur sed, eleifend ac, lectus. Nulla facilisi. Pellentesque eget lectus. Proin eu metus. Sed porttitor. In hac habitasse platea dictumst. Suspendisse eu lectus. Ut mi mi, lacinia sit amet, placerat et, mollis vitae, dui. Sed ante tellus, tristique ut, iaculis eu, malesuada ac, dui. Mauris nibh leo, facilisis non, adipiscing quis, ultrices a, dui.

Morbi luctus, wisi viverra faucibus pretium, nibh est placerat odio, nec commodo wisi enim eget quam. Quisque libero justo, consectetur a, feugiat vitae, porttitor eu, libero. Suspendisse sed mauris vitae elit sollicitudin malesuada. Maecenas ultricies eros sit amet ante. Ut venenatis velit. Maecenas sed mi eget dui varius euismod. Phasellus aliquet volutpat odio. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Pellentesque sit amet pede ac sem eleifend consectetur. Nullam elementum, urna vel imperdiet sodales, elit ipsum pharetra ligula, ac pretium ante justo a nulla. Curabitur tristique arcu eu metus. Vestibulum lectus. Proin mauris. Proin eu nunc eu urna hendrerit faucibus. Aliquam auctor, pede consequat laoreet varius, eros tellus scelerisque quam, pellentesque hendrerit ipsum dolor sed augue. Nulla nec lacus.

Suspendisse vitae elit. Aliquam arcu neque, ornare in, ullamcorper quis, commodo eu, libero. Fusce sagittis erat at erat tristique mollis. Maecenas sapien libero, molestie et, lobortis in, sodales eget, dui. Morbi ultrices rutrum lorem. Nam elementum ullamcorper leo. Morbi dui. Aliquam sagittis. Nunc placerat. Pellentesque tristique sodales est. Maecenas imperdiet lacinia velit. Cras non urna. Morbi eros pede, suscipit ac, varius vel, egestas non, eros. Praesent malesuada, diam id pretium elementum, eros sem dictum tortor, vel consectetur odio sem sed wisi.

4 EXPERIMENTAL SYSTEM: TASKS AND THE BLACKBOARD GAME

Our experiment separates task construction from collective decision-making. We first create a team-allocation problem with a known answer and divide its evidence among language-model agents. We then let the agents exchange that evidence and revise their answers through a shared message board. A controller can post on the same board to encourage a chosen answer. The resulting game lets us measure how communication changes both the information available to agents and their collective decision.

4.1 CONSTRUCTING THE TASK

A team-allocation problem. Our tasks adapt the Team Allocation setting of MuSR (Sprague et al., 2024), which tests multistep reasoning from natural-language evidence. Three people must be assigned to two jobs: one person works alone on the first job, and the remaining pair works together on the second. There are therefore three possible answers. For example, suppose the jobs are to analyze field measurements and organize a community workshop. The choices are:

Analyze measurements	Organize the workshop
Alice	Bruno and Chandra
Bruno	Alice and Chandra
Chandra	Alice and Bruno

Choosing among them requires evidence about individual skills and cooperation within each possible pair. We generate our own instances of this task family.

Establishing the answer before writing the evidence. We first construct a hidden numerical model with six skill values (one per person and job) and three pairwise cooperation values. Each value lies in $\{1, 2, 3\}$. If person p takes the first job and people u, v take the second, our allocation score is

$$S(p; u, v) = s_{p,1} + s_{u,2} + s_{v,2} + c_{u,v}, \quad (2)$$

where $s_{p,t}$ measures person p 's skill for job t and $c_{u,v}$ measures cooperation. We evaluate all three allocations and retain worlds with a unique highest-scoring answer. Thus, the correct answer is determined by code before any language-model interaction. Agents receive the scenario and evidence, while the numerical model and scores remain hidden.

Turning the hidden model into evidence cards. Following MuSR's approach of constructing a reasoning problem before rendering it in language (Sprague et al., 2024), we create short accounts of behavior that support facts about skills or cooperation. For example, an account of Alice identifying and repairing a spreadsheet error provides evidence of her analytical ability. We call each short account an *evidence card*. Every card is linked to a proposition that can be checked against the hidden model, such as a person's skill level or a comparison between two people's skills. An LLM writes the accounts; equality comparisons use fixed templates. We check that the text supports its intended proposition and does not disclose the hidden scores or recommend an answer. This preserves the distinction between the exact answer, established by the numerical model, and the natural-language reasoning needed to infer it from the cards.

Distributing incomplete information. We assign one private card to each of $N = 24$ agents. A single card leaves several answers plausible, whereas combining cards provides stronger evidence. For example, one agent may know about Alice's analytical ability and another about Bruno and Chandra's cooperation. We assess this partial-information ambiguity by enumerating the numerical models consistent with each private card and checking which allocation wins in each model. The assignment includes complementary cards that favor the correct answer when considered together. Agents may initially guess correctly, but they must exchange evidence to access information beyond their own card.

We freeze the task and evidence before play. Each agent first votes using only its private card. These initial votes are saved and reused across matched experimental conditions. For each model call, we shuffle the A/B/C labels assigned to the three allocations and translate the response back to the chosen allocation, preventing a fixed answer letter from becoming a shared social cue.

4.2 PLAYING THE BLACKBOARD GAME

The *blackboard* is a shared collection of short-lived messages. Agents take turns reading a small random sample of posts, choosing an answer, and optionally adding a post of their own. A message posted early can therefore influence later decisions within the same round. For example, an agent may post the card about Alice's analytical ability; another agent who reads it can combine it with its own cooperation evidence and reconsider the allocation. Access to the board is limited: agents see only the messages they sample.

Agent state and memory. Agent i maintains a current answer X_i , an active evidence set K_i , and a historical set H_i containing every card it has received. Only active evidence is available in its prompt. At each round boundary, each active card survives independently with probability ρ ; historical records are retained. Reading a previously lost card makes it active again. The parameter ρ therefore controls how long communicated evidence remains usable. Previous private explanations are recorded for analysis but are not fed back to agents.

One round. Each round has a boundary stage, a controller posting stage, and N sequential agent updates. The controller first samples votes and decides whether to act. The previous round’s messages then expire, evidence persistence is applied, and an active controller posts its messages. During each subsequent update:

1. One focal agent is sampled uniformly from the population. It reads up to q live board messages, sampled uniformly without replacement, excluding its own posts.
2. The agent receives the question, answer options, active evidence, previous vote, and sampled messages. One LLM call returns an updated vote, a private explanation, and an optional public message.
3. The vote is updated immediately, attached evidence from the sampled reports is added to the agent’s active and historical sets, and any new public message becomes available to later updates in the same round.

Focal agents are sampled independently at each position, so a round need not update every agent once. Messages live only for their creation round; an empty board supplies no social messages. Episodes run for a fixed number of rounds, including after consensus.

Messages and evidence transmission. An ordinary agent may post a `REPORT`, ask for evidence with a `REQUEST`, or post nothing. A report contains public text and may attach one evidence identifier available to its author. Requests carry no evidence. Posts also carry the author’s current allocation choice, and replies may refer only to messages visible in that update. The prompt distinguishes a participant’s interpretation from the canonical evidence text attached to a report. Only reading a report with a valid attachment changes the exact evidence state. Posting alone does not transmit a card to anyone, and public prose is not counted as verified evidence. This separates the transmission of task evidence from persuasion through an agent’s interpretation of that evidence.

Feedback through the same public board. Let Z be the controller’s target, which may be correct or incorrect, and let

$$n_k = \sum_{i=1}^N \mathbb{I}[X_i(k) = Z], \quad x_k = n_k/N \quad (3)$$

be its support before round k . The controller samples q_c distinct agents and observes their votes, giving a target-support count Y_k . The binary feedback policy in Section 5 chooses between silence ($U_k = 0$) and communication ($U_k = 1$); lower sampled target support makes communication more likely. This decision uses votes only.

The controller has a fixed pool of true evidence cards. To favor an incorrect answer, it selectively discloses evidence that keeps that answer plausible while withholding evidence that would weaken it. Its reports remain truthful at the fact level.

On active rounds, an adaptive communication policy chooses to publish up to b evidence reports, make one request for evidence, or post one `DIRECTIVE` directing attention to a comparison. The latter two modes carry no evidence. An LLM selects the mode and eligible report identifiers using sampled votes, previous public messages, the allowed fact pool, and posting history. It has no access to agents’ private evidence or explanations. The selected reports reproduce the frozen evidence text, while requests and directives use fixed templates. This communication choice occurs only after the binary decision to act. A report-only variant instead posts exactly b distinct evidence reports on each active round.

The controller appears under a stable participant identity. Its messages enter the same board and have no sampling priority. Consequently, b bounds posting capacity, while actual exposure depends on which messages agents read. We record posts, reads, evidence acquisition and reactivation, and vote changes separately. Tracking support for both Z and the correct answer, alongside active and historical evidence, allows us to distinguish successful steering from improved task performance.

5 EFFICIENCY OF FEEDBACK CONTROL

We use a solvable blackboard reference to separate sensing, message exposure, and behavioral response. Let $n_k \in \{0, \dots, N\}$ count supporters of the controller target Z before round k , and write

$x_k = n_k/N$. One feedback cycle has causal order $n_k \rightarrow Y_k \rightarrow U_k \rightarrow n_{k+1}$, where $U_k = 0$ denotes silence and $U_k = 1$ active posting. Both branches contain M population updates, usually $M = N$. Posting b messages changes the chance of exposure; it does not guarantee b controlled revisions. The results below are exact for the reference kernels, rather than a fitted description or an energetic thermodynamics of the full language-model system.

5.1 SENSING, POSTING, AND SAMPLED EXPOSURE

The controller senses q_c agents without replacement. For $Y = y$ sampled target supporters, the sensor and soft feedback policy are

$$S(y | n) = \frac{\binom{n}{y} \binom{N-n}{q_c-y}}{\binom{N}{q_c}}, \quad (4)$$

$$\pi(1 | y) = \sigma \left[\beta \left(\theta - \frac{y}{q_c} \right) \right], \quad \sigma(z) = (1 + e^{-z})^{-1}, \quad (5)$$

$$a_n = P(U = 1 | n) = \sum_y S(y | n) \pi(1 | y). \quad (6)$$

Suppose the frozen eligible board contains B peer messages and, when active, b controller messages. Each focal update samples $s = \min(q, B + b)$ messages uniformly without replacement. Its probability of seeing at least one controller message is

$$w_b = 1 - \frac{\binom{B}{s}}{\binom{B+b}{s}}. \quad (7)$$

Here $\binom{B}{s} = 0$ when $s > B$. The theory uses the number of messages *actually posted*; a nominal budget must be replaced by the realized count when fewer messages are posted. Thus q_c, θ, β govern the posting decision, while B, b, q govern exposure. For positive q , exposure saturates as b grows.

5.2 FROM ONE AGENT UPDATE TO A FULL FEEDBACK ROUND

One update: willingness to revise and preference after revision. At each microscopic update, one agent is selected uniformly from the N agents. We describe its response using two probabilities. The parameter $\gamma \in [0, 1]$, called *kinetic compliance*, is the probability that the agent revises its vote. Conditional on revising, $p \in [0, 1]$ is the probability that it selects the controller target Z ; with probability $1 - p$ it selects a non-target alternative. With probability $1 - \gamma$, it keeps its current vote. Revision here means making a new choice, which can agree with the old one. Thus γ is not the probability of an observed vote change. For example, $\gamma = 0.4$ and $p = 0.8$ mean a 40% chance of making a new choice and, conditional on that choice, an 80% chance of selecting Z .

A kernel is a table of transition probabilities. Write $K_{\gamma,p}(m | n)$ for the probability that this single update leaves m target supporters, given n initially. The symbol K denotes this transition table, or *kernel*; the subscripts specify which revision probabilities it uses. Only one agent updates, so the count can increase by one, decrease by one, or stay fixed:

$$K_{\gamma,p}(n+1 | n) = \gamma(1 - n/N)p, \quad (8)$$

$$K_{\gamma,p}(n-1 | n) = \gamma(n/N)(1-p), \quad (9)$$

$$K_{\gamma,p}(n | n) = 1 - K_{\gamma,p}(n+1 | n) - K_{\gamma,p}(n-1 | n). \quad (10)$$

For the first line, the selected agent must be a non-supporter, an event of probability $1 - n/N$, then revise with probability γ , and choose Z with probability p . The second line reverses this reasoning: a supporter is selected, revises, and chooses a non-target alternative. The last line collects all ways to leave the supporter count unchanged.

Two response classes: without and with exposure. An agent may respond differently when its sampled messages include controller content. We therefore use the same transition rule with two parameter pairs:

$$K_0 = K_{\gamma_0, p_0}, \quad K_1 = K_{\gamma_1, p_1}.$$

The subscript 0 means *no controller message was sampled*: γ_0 is the baseline revision probability and p_0 is the baseline probability of choosing Z after revision. The subscript 1 means *at least one controller message was sampled*: γ_1 and p_1 are the corresponding exposed-response probabilities. These are effective behavioral parameters, not message counts. Exposure may change willingness to revise, preference for the target, or both. In particular, p_0 need not be zero: peer messages can support Z without controller exposure.

An active controller does not expose every update. When the controller posts b messages, a focal update sees controller content with probability w_b from Eq. equation 7. It therefore follows K_1 with probability w_b and K_0 otherwise. The resulting single-update kernel is their probability-weighted mixture,

$$K_b = (1 - w_b)K_0 + w_bK_1 = K_{\gamma_b, p_b}. \quad (11)$$

Here the subscript b refers to the number of posted messages, whereas 0 and 1 above label exposure classes. The mixture has the same form as the original rule, with effective parameters

$$\gamma_b = (1 - w_b)\gamma_0 + w_b\gamma_1, \quad (12)$$

$$p_b = \frac{(1 - w_b)\gamma_0 p_0 + w_b\gamma_1 p_1}{\gamma_b}. \quad (13)$$

The first expression is the overall probability of revision. The numerator of the second is the probability of revising and choosing Z ; dividing by γ_b gives the preference conditional on revision. Thus p_b weights each exposure class by how often it produces revision, not just how often it is encountered. If $\gamma_b = 0$, no revision occurs, K_b leaves every state unchanged, and p_b is unnecessary.

From one update to a round. The kernels K describe one agent update. We use Q for the transition table of a full round containing M successive updates. Repeating the same update rule M times gives a matrix power:

$$Q_0 = K_0^M, \quad Q_1 = K_b^M. \quad (14)$$

For Q , the subscript labels the *controller action*: Q_0 is the silent-round kernel and Q_1 the active-round kernel. An active round uses K_b , not K_1 , because individual updates may still miss the controller messages. A silent round uses K_0 throughout, but can still change votes through baseline social interaction. Finally, before the controller action is known, the round kernel averages over its posting probability a_n :

$$Q_\pi(m | n) = (1 - a_n)Q_0(m | n) + a_nQ_1(m | n). \quad (15)$$

Thus the notation follows three steps: K_0, K_1 specify responses to exposure; K_b averages over whether exposure occurs; and Q_0, Q_1, Q_π describe full rounds, first conditional on the controller action and then averaged over the policy. The matrix-power construction assumes a frozen board, common response parameters across agents and updates, and the same unexposed baseline in both controller branches.

5.3 RAW AND AVAILABILITY-NORMALIZED SUSCEPTIBILITY

At a common initial state, the primary response observable is

$$\chi_b(x) = \mathbb{E}[x' | U = 1, x] - \mathbb{E}[x' | U = 0, x]. \quad (16)$$

The one-update mean obeys $\mathbb{E}[x_{j+1} | x_j] = (1 - \gamma/N)x_j + \gamma p/N$. Consequently the exact round mean and susceptibility are

$$F_M(x; \gamma, p) = p + (x - p)(1 - \gamma/N)^M, \quad (17)$$

$$\begin{aligned} \chi_b(x) &= F_M(x; \gamma_b, p_b) - F_M(x; \gamma_0, p_0) \\ &= p_b(1 - r_b) - p_0(1 - r_0) + (r_b - r_0)x, \end{aligned} \quad (18)$$

where $r_b = (1 - \gamma_b/N)^M$ and $r_0 = (1 - \gamma_0/N)^M$. The reference predicts an affine dependence on initial target share, with budget entering through exposure. If $r_b \neq r_0$, the zero-response point is $x_b^* = [p_0(1 - r_0) - p_b(1 - r_b)]/(r_b - r_0)$; it need not lie in $[0, 1]$. When $\gamma_0 = \gamma_1 = \gamma$, the state dependence cancels:

$$\chi_b(x) = w_b(p_1 - p_0) [1 - (1 - \gamma/N)^M]. \quad (19)$$

Exposure and compliance therefore impose distinct limits: being seen more often and changing behavior after exposure are separate mechanisms.

To express the same branch contrast relative to the remaining non-target population, define

$$\chi_{b,\text{avail}}(x) = \frac{\chi_b(x)}{1-x}, \quad x < 1. \quad (20)$$

The saturated state $x = 1$ is excluded. This diagnostic can be negative or exceed one: the contrast includes differential retention of existing target supporters as well as additional adoption. It is neither a conversion probability nor a bounded efficiency. For occupancy $p(n)$, the corresponding available-mass-weighted summary is

$$\bar{\chi}_{\text{avail}}[p] = \frac{\sum_{n=0}^{N-1} p(n) \chi_b(n/N)}{\sum_{n=0}^{N-1} p(n) (1 - n/N)}, \quad (21)$$

when the denominator is positive. Both sums exclude saturation; this ratio avoids an unweighted average of potentially large near-saturation ratios.

5.4 ACTION INFORMATION AND INFORMATION-RESPONSE EFFICIENCY

The exact state-local action-to-population information is the weighted Jensen-Shannon divergence of the two evolving round kernels:

$$T_\pi(n) = I(U; n' | n) = \sum_u P(u | n) \sum_m Q_u(m | n) \log_2 \frac{Q_u(m | n)}{Q_\pi(m | n)}. \quad (22)$$

With $H_2(a) = -a \log_2 a - (1-a) \log_2 (1-a)$, $0 \leq T_\pi(n) \leq H_2(a_n)$. Thus the action-information fraction is

$$\eta_{\text{IF}}(n) = \frac{T_\pi(n)}{H_2(a_n)} \in [0, 1], \quad (23)$$

when $H_2(a_n) > 0$. It measures how much conditional action uncertainty is visible in the next state, rather than directional motion.

Because $x' \in [0, 1]$, Pinsker's inequality gives

$$T_\pi(n) \geq B_{\text{IR}}(n) \equiv \frac{2a_n(1-a_n)}{\ln 2} \chi_b(n/N)^2. \quad (24)$$

The bounded information-response efficiency is therefore

$$\eta_{\text{IR}}(n) = \frac{2a_n(1-a_n)\chi_b(n/N)^2}{(\ln 2)T_\pi(n)} \leq 1, \quad (25)$$

for $T_\pi(n) > 0$. This bound uses the *raw* susceptibility, not its availability-normalized counterpart, and does not require Q_0 to be the identity. An informative action channel can have little mean response if its effect is mainly on fluctuations or higher moments.

5.5 EXCESS CURRENT AND FINITE-TIME PATH ACCOUNTING

Let $p_k(n)$ denote the pre-round ensemble and $\mu_u(n) = \mathbb{E}[n' - n | n, U = u]$. Controller-excess current subtracts the silent counterfactual:

$$J_{c,k}^{\text{ex}} = \sum_n p_k(n) a_n [\mu_1(n) - \mu_0(n)] = N \sum_n p_k(n) a_n \chi_b(n/N). \quad (26)$$

Equivalently, $J_{c,k}^{\text{ex}} = J_{\text{tot},k} - J_{\text{sil},k}$, where $J_{\text{tot},k} = \sum_n p_k(n) [(1-a_n)\mu_0(n) + a_n\mu_1(n)]$ and $J_{\text{sil},k} = \sum_n p_k(n) \mu_0(n)$.

For interior preferences, define $h_0 = \ln[p_0/(1-p_0)]$ and $h_b = \ln[p_b/(1-p_b)]$, with $h_u = h_0$ under silence and $h_u = h_b$ under active posting. On supported transitions, weighted reversibility gives

$$\ln \frac{Q_u(m | n)}{Q_u(n | m)} = \ln \binom{N}{m} - \ln \binom{N}{n} + h_u(m - n). \quad (27)$$

Let $p_{Y,k}(y) = \sum_n p_k(n)S(y | n)$ and let p_{k+1} be the ensemble propagated by Q_π . The forward and reverse-reference path measures are

$$P_F(n, y, u, m) = p_k(n)S(y | n)\pi(u | y)Q_u(m | n), \quad (28)$$

$$P_R(m, y, u, n) = p_{k+1}(m)p_{Y,k}(y)\pi(u | y)Q_u(n | m). \quad (29)$$

Their log ratio yields the exact finite-time identity

$$\Sigma_k = D_{\text{KL}}(P_F \| P_R) = \Delta S_{\text{sys},k} + I(n_k; Y_k) + \mathcal{A}_k \geq 0, \quad (30)$$

where all terms use natural logarithms and

$$S_{\text{sys}}[p] = -\sum_n p(n) \ln p(n) + \sum_n p(n) \ln \binom{N}{n}, \quad (31)$$

$$\mathcal{A}_k = \sum_n p_k(n) [(1 - a_n)h_0\mu_0(n) + a_n h_b \mu_1(n)]. \quad (32)$$

The entropy change is retained; stationarity is not assumed.

For $\mathcal{A}_k \geq 0$ and positive denominator, the bounded *total path-transport fraction* is

$$\eta_{\text{path},k} = \frac{\mathcal{A}_k}{\Sigma_k - \Delta S_{\text{sys},k}} = \frac{\mathcal{A}_k}{\mathcal{A}_k + I(n_k; Y_k)} \in [0, 1]. \quad (33)$$

It includes autonomous baseline transport and is not a controller-only thermodynamic efficiency. Indeed, with $J_{\text{adv},k} = \sum_n p_k(n)a_n\mu_1(n)$,

$$\mathcal{A}_k - h_0 J_{\text{sil},k} = h_0 J_{c,k}^{\text{ex}} + (h_b - h_0)J_{\text{adv},k}. \quad (34)$$

This excess-affinity diagnostic need not be nonnegative and is not itself a second-law-bounded expenditure. A sufficient attribution condition is a neutral baseline affinity, $p_0 = 1/2$ or $h_0 = 0$. Only under this additional condition does the path fraction become

$$\eta_{\text{th},k}^{(h_0=0)} = \frac{h_b J_{\text{adv},k}}{h_b J_{\text{adv},k} + I(n_k; Y_k)} \in [0, 1], \quad (35)$$

provided $h_b J_{\text{adv},k} \geq 0$ and the denominator is positive. Neutral affinity does not imply zero silent drift. In general, the pair $(J_{c,k}^{\text{ex}}, \Sigma_k)$ describes excess steering versus total coarse path irreversibility without assigning baseline transport to the controller.

Empirical scope. Exposure frequency and exposed/unexposed transition rates can test the reference. If $p_{+,r}$ and $p_{-,r}$ are conditional target-entry and exit probabilities in response class r , then $\gamma_r = p_{+,r} + p_{-,r}$ and $p_r = p_{+,r}/(p_{+,r} + p_{-,r})$ when the denominator is nonzero. Shared parameters should predict held-out budgets and states; observational exposure classes alone do not establish a causal effect. Paired active and silent continuations should match the complete state, including board contents and epistemic information. Board growth, self-exclusion, message content, and heterogeneous epistemic state can violate the frozen, homogeneous reference. Such departures call for explicit conditioning or time-dependent kernels. The safe bounded information quantities are η_{IF} and η_{IR} ; controller-specific thermodynamic attribution requires additional assumptions.

6 RESULTS

We organize the results around three complementary questions. First, how much can the controller move the population? Second, how efficiently is the controller-dependent information expressed as coherent population motion? Third, how efficiently is the finite-time thermodynamic control budget converted into directed current? These questions correspond respectively to the susceptibility χ , the information–response efficiency η_{IR} , and the thermodynamic efficiency η_{th} .

6.1 CONTROL SUSCEPTIBILITY IS CAUSAL, FINITE, AND STATE DEPENDENT

We first measure the response to advocacy under matched initial conditions. For a fixed initial target fraction $x_0 = n_Z/N$, the susceptibility is

$$\chi(x_0) = \mathbb{E}[\Delta x | \text{ADV}, x_0] - \mathbb{E}[\Delta x | \text{NOOP}, x_0].$$




Figure 1: **Empirical and theoretical susceptibility.** (a) State-matched ADVOCATE-minus-NoOp response with 95% bootstrap intervals and the finite-compliance calibration from microscopic transition rates. (b) Exact theory showing how kinetic compliance γ changes the response amplitude without moving the intrinsic set point.

Holding the epistemic composition fixed, advocacy produces a positive shift toward the controller target at every tested state:

$$\chi(0.25) = 0.163, \quad \chi(0.375) = 0.100, \quad \chi(0.50) = 0.104.$$

The corresponding 95% matched-pair bootstrap intervals are $[0.119, 0.202]$, $[0.060, 0.144]$, and $[0.060, 0.148]$ (Fig. 1a). Thus the controller has a measurable causal influence even when the initial knowledge distribution is held fixed. The response is largest at the lowest tested target fraction and does not increase with additional pre-existing support over the tested range.

Microscopic controlled transitions provide a direct calibration of the finite-compliance description. Pooling the state-matched controlled updates gives an effective directional affinity $h_{\text{eff}} = 3.83$ and kinetic compliance $\gamma_{\text{eff}} = 0.372$. Substituting these into

$$\chi_{h,\gamma}(x) = [\sigma(h) - x] \left[1 - \left(1 - \frac{\gamma}{N} \right)^b \right]$$

gives the dashed finite-compliance reference in Fig. 1a. Because h_{eff} and γ_{eff} are estimated from the same microscopic trajectories, this is a calibrated reference rather than a parameter-free prediction. It nevertheless shows that separating directional bias from kinetic compliance removes much of the discrepancy with the idealized copying limit.

The exact theory also clarifies the role of compliance itself (Fig. 1b). Changing γ rescales the magnitude of the response while leaving the zero-response set point $x^* = \sigma(h)$ unchanged. This separation is useful empirically: h controls which direction is favored, whereas γ controls how often a controlled opportunity actually produces a revision.

6.2 STATE-LOCAL DYNAMICS REVEAL WHERE CONTROL IS EXPRESSED

The longer stochastic-feedback experiment gives access to the trajectory-level state dependence of control. Figure 2a shows the mean signed round-to-round drift $\mathbb{E}[\Delta x \mid x, b]$ for $q_c = 12$. Because the controller advocates an incorrect target, the uncontrolled epistemic dynamics usually drive x downward. Increasing b makes this correction progressively weaker, particularly at low and intermediate target support. This quantity is descriptive rather than causal: the matched experiment above remains the clean measurement of χ .

The controller action also leaves a strongly state-dependent information signature. We measure the reduced-state conditional mutual information

$$T_\pi(n) \equiv I(U_k; n_{k+1} \mid n_k = n),$$

the action-to-population information induced by the chosen feedback policy π . Here U_k is the policy-selected controller action and n_{k+1} is the population state after actuation. The empirical excess above a within-state permutation null is shown in Fig. 2b. At small b the signal is weak, whereas at $b = 18$ a clear interior band appears. The exact finite-state theory produces the same basic geometric mechanism (Fig. 2c): local information is high only where the action is not already




Figure 2: **State-local response and information.** (a) Empirical signed population drift as a function of target support and actuation budget at $q_c = 12$. (b) Empirical state-local action-to-population information above a within-state permutation null. (c) Exact theoretical state-local information landscape for the baseline single-affinity finite-compliance feedback controller. (d) Number of observed rounds supporting each empirical state bin.

deterministic and where the advocacy and NoOp transition kernels are sufficiently separated. The support map in Fig. 2d is included because the higher state bins are visited less often and should not be interpreted with the same confidence as densely sampled cells.

The robust conclusion is that control information is not uniform over population state space and stronger actuation creates an informative interior region. We do not interpret the precise ridge location as a phase boundary because n_Z is only a coarse-grained state.

6.3 FEEDBACK GAIN CONCENTRATES RATHER THAN SIMPLY INCREASES INFORMATION

The theory makes a more specific prediction about the feedback policy. The gain β controls how sharply the controller switches between advocacy and NoOp after sensing. At small β , the action remains uncertain over a wide range of states; at large β , the policy becomes nearly deterministic except close to its switching threshold. Consequently increasing β does not monotonically increase $T_{\pi}(n)$. Instead, the informative band narrows and moves toward the decision boundary (Fig. 3a).

The state-averaged theory curves make the same point in one dimension (Fig. 3b): for fixed actuation budget, larger β can reduce the average action-to-population information even though the controller is more decisive. This separates controller *decisiveness* from controller *information transmission* and suggests a focused validation rather than a broad search: $\beta \in \{1, 4, 8\}$ is sufficient to test whether the empirical information band sharpens as predicted.

6.4 INFORMATION-RESPONSE EFFICIENCY INCREASES WITH ACTUATION

The information-response efficiency asks how much of the controller-dependent change in the full next-state distribution is expressed as a coherent shift of the mean target fraction. Using the bounded quantity defined in Section 5, we report the finite-horizon information-response efficiency $\eta_{\text{IR}}^{(H)}$, obtained as the ratio of the accumulated occupancy-weighted Pinsker response terms to the accumulated state-local action-to-population information over the evaluation horizon. The empirical values are small in absolute terms but systematic across the resource grid (Fig. 4a). They range from




Figure 3: **Exact theory for feedback gain.** (a) State-local action-to-population information for increasing policy gain β . A sharper policy concentrates information around the switching region and suppresses it where the action saturates. (b) State-averaged local information versus actuation budget for the theory parameter sweeps; the β family shows that a more deterministic policy need not transmit more action information overall.

approximately 0.4% to 2.5% of the rigorous Pinsker bound and increase with b for every sensing budget. At $q_c = 12$, for example,

$$\eta_{\text{IR}} = 0.0043, 0.0135, 0.0228 \quad \text{for } b = 6, 12, 18.$$

Thus stronger actuation does more than increase the size of the population response: within this operating regime it also makes a larger fraction of the action-dependent distributional change visible as coherent mean motion. The small absolute values should not be interpreted as energetic losses; the Pinsker inequality supplies a conservative information-theoretic ceiling and need not be tight for these transition distributions.

6.5 THERMODYNAMIC EFFICIENCY EXPOSES THE SENSING–ACTUATION TRADE-OFF

The single-affinity path identity leads to a bounded thermodynamic control efficiency without introducing an opposing load. Removing the transient system entropy term from the finite-time path irreversibility gives

$$\Sigma - \Delta S_{\text{sys}} = hJ_c + I(n_k; Y_k),$$

and therefore

$$\eta_{\text{th}} = \frac{hJ_c}{\Sigma - \Delta S_{\text{sys}}} = \frac{hJ_c}{hJ_c + I(n_k; Y_k)} \leq 1.$$

The numerator is the affinity-weighted target-directed current; the denominator is the non-storage part of the finite-time thermodynamic accounting. Thus η_{th} asks what fraction of the sensing-plus-actuation expenditure is expressed as directed population motion. Appendix A.2 gives the complete path-measure derivation.

For empirical trajectories, we report the finite-horizon thermodynamic efficiency $\eta_{\text{th}}^{(H)}$. The affinity-weighted controlled current and scalar target-sensing information are accumulated cycle by cycle, and the efficiency ratio is formed only after this accumulation. We therefore do not average the cycle-wise efficiencies, and we do not compute sensing information from a single pooled occupancy distribution.

Figure 4b shows the exact single-affinity prediction for the same sensing budgets used experimentally, with $h = 2$, $\gamma = 0.75$, $\beta = 4$, $\theta = 1/2$, and $p_0(n) = \text{Binomial}(24, 0.25)$. At fixed q_c , the efficiency rises with actuation because the controlled current grows while the sensing information is fixed by the sensor channel. At fixed b , larger sensing budgets lower the efficiency unless the additional information produces a commensurate increase in current. For $b = 18$, for example,

$$\eta_{\text{th}} = 0.985, 0.964, 0.930 \quad \text{for } q_c = 6, 12, 18.$$

The result is therefore not that sensing is intrinsically undesirable, but that information that does not change the action-relevant population response is a thermodynamic overhead in this control objective.




Figure 4: **The two control efficiencies and their resource dependence.** (a) Empirical information-response efficiency across the sensing-actuation grid. (b) Exact single-affinity thermodynamic efficiency $\eta_{th} = hJ_c / (hJ_c + I_{sens})$ as the actuation budget is varied for the three sensing budgets used experimentally. (c) Empirical sensing information, which rises strongly with q_c without a proportional increase in controlled current.

The empirical sensing measurements support precisely this mechanism (Fig. 4c). Increasing q_c substantially raises $I(n_k; Y_k)$, whereas the calibrated controlled current changes much less across the same resource grid. Combined with the close agreement between the measured current and the finite-compliance current reported above, this yields a direct design prediction: in the present regime, adding actuation improves both η_{IR} and η_{th} , whereas additional sensing is efficient only when the policy can convert it into additional directed current.

Taken together, the results support a single resource-conversion picture of feedback control. Susceptibility identifies where the population can be moved; state-local information identifies where the action is dynamically visible; η_{IR} quantifies how coherently that information appears as mean motion; and η_{th} quantifies how efficiently the finite-time sensing-actuation budget is converted into affinity-weighted directed current. In the present regime, actuation is the dominant useful resource, while sensing displays clear diminishing returns. The theory additionally predicts that feedback gain β should reshape, rather than simply amplify, the information channel, providing a focused next validation of the framework.

7 CONCLUSION

Curabitur tellus magna, porttitor a, commodo a, commodo in, tortor. Donec interdum. Praesent scelerisque. Maecenas posuere sodales odio. Vivamus metus lacus, varius quis, imperdiet quis, rhoncus a, turpis. Etiam ligula arcu, elementum a, venenatis quis, sollicitudin sed, metus. Donec nunc pede, tincidunt in, venenatis vitae, faucibus vel, nibh. Pellentesque wisi. Nullam malesuada. Morbi ut tellus ut pede tincidunt porta. Lorem ipsum dolor sit amet, consectetur adipiscing elit. Etiam congue neque id dolor.

Donec et nisl at wisi luctus bibendum. Nam interdum tellus ac libero. Sed sem justo, laoreet vitae, fringilla at, adipiscing ut, nibh. Maecenas non sem quis tortor eleifend fermentum. Etiam id tortor ac mauris porta vulputate. Integer porta neque vitae massa. Maecenas tempus libero a libero posuere dictum. Vestibulum ante ipsum primis in faucibus orci luctus et ultrices posuere cubilia Curae; Aenean quis mauris sed elit commodo placerat. Class aptent taciti sociosqu ad litora torquent per conubia nostra, per inceptos hymenaeos. Vivamus rhoncus tincidunt libero. Etiam elementum pretium justo. Vivamus est. Morbi a tellus eget pede tristique commodo. Nulla nisl. Vestibulum sed nisl eu sapien cursus rutrum.

REFERENCES

- Anthropic. Patterns and problems in emerging multiagent systems, August 2026. URL <https://www.anthropic.com/research/multiagent-systems>.
- Ariel Flint Ashery, Luca Maria Aiello, and Andrea Baronchelli. Emergent social conventions and collective bias in LLM populations. *Science Advances*, 11(20):eadu9368, 2025. doi: 10.1126/sciadv.adu9368. URL <https://doi.org/10.1126/sciadv.adu9368>.

- Claudio Castellano, Santo Fortunato, and Vittorio Loreto. Statistical physics of social dynamics. *Reviews of Modern Physics*, 81(2):591–646, 2009. doi: 10.1103/RevModPhys.81.591. URL <https://doi.org/10.1103/RevModPhys.81.591>.
- Ayush Chopra. Intelligence is in the interaction. MIT Media Lab, February 2026. URL <https://www.media.mit.edu/articles/intelligence-is-in-the-interaction/>.
- Batu El, Jinhee Paeng, Fatih Dinc, Shiye Su, Mete Erdogan, Aneesh Pappu, Haotian Ye, Wanjia Zhao, Surya Ganguli, and James Zou. Physics of agents: Statistical mechanics predicts collective behavior of ai agents, 2026. URL <https://arxiv.org/abs/2608.16578>.
- Google DeepMind, Schmidt Sciences, Cooperative AI Foundation, Advanced Research and Invention Agency, and Google.org. Investing in multi-agent ai safety research, June 2026. URL <https://deepmind.google/blog/investing-in-multi-agent-ai-safety-research/>.
- Ryan Greenblatt, Ajeya Cotra, and Hjalmar Wijk. Brief independent investigation of agents’ behavior, reasoning and collaboration in the openai / hugging face hacking incident. METR and Redwood Research, August 2026. URL <https://metr.org/blog/2026-08-26-openai-hugging-face-incident-investigation/>.
- Yukun Jiang, Yage Zhang, Xinyue Shen, Michael Backes, and Yang Zhang. “Humans welcome to observe”: A First Look at the Agent Social Network Moltbook. *CoRR*, abs/2602.10127, 2026. URL <https://arxiv.org/abs/2602.10127>.
- Alexander Meulemans, Maciej Wołczyk, Marissa A. Weis, Rajai Nasser, Roberta Rocca, Seijin Kobayashi, Guillaume Lajoie, Angelika Steger, Blake Richards, Marcus Hutter, James Manyika, Rif A. Saurous, João Sacramento, and Blaise Agüera y Arcas. A game theory for foundation models shows new paths to rational cooperation through similarity inference, 2026. URL <https://arxiv.org/abs/2608.03958>.
- OpenAI. On the navier–stokes millennium prize problem, September 2026. URL <https://openai.com/index/navier-stokes-solution/>.
- Christoph Riedl. Emergent coordination in multi-agent language models. In *International Conference on Learning Representations*, 2026. URL <https://arxiv.org/abs/2510.05174>.
- Zayne Sprague, Xi Ye, Kaj Bostrom, Swarat Chaudhuri, and Greg Durrett. MuSR: Testing the limits of chain-of-thought with multistep soft reasoning. In *International Conference on Learning Representations*, 2024. URL https://proceedings.iclr.cc/paper_files/paper/2024/hash/3f8c7eb848ffec848f3ed2b7ca44915d-Abstract-Conference.html.

A THEORY AND DERIVATIONS

A.1 DERIVATION OF THE INFORMATION–RESPONSE BOUND

This appendix derives the bound

$$T_{\pi}(n) \geq \frac{2a_n(1-a_n)}{\ln 2} \chi_{h,\gamma}(n/N)^2 \quad (36)$$

used to define the information–response efficiency. The derivation proceeds in three steps: (i) expressing the conditional mutual information as a weighted Jensen–Shannon divergence, (ii) applying Pinsker’s inequality, and (iii) bounding the change in the mean target fraction by total variation.

Throughout the derivation, the current population state n is held fixed.

A.1.1 CONDITIONAL MUTUAL INFORMATION AS A JENSEN–SHANNON DIVERGENCE

Let the binary control variable be

$$U \in \{0, 1\},$$

where $U = 0$ denotes no intervention and $U = 1$ denotes advocacy of the controller target Z . At population state n , define

$$\Pr(U = 1 \mid n) = a_n, \quad \Pr(U = 0 \mid n) = 1 - a_n. \quad (37)$$

The corresponding next-round target-count distributions are

$$Q_0(n' \mid n) = \Pr(n' \mid U = 0, n), \quad (38)$$

$$Q_1(n' \mid n) = \Pr(n' \mid U = 1, n). \quad (39)$$

For notational convenience, while n is fixed, write

$$P(n') \equiv Q_0(n' \mid n), \quad Q(n') \equiv Q_1(n' \mid n), \quad a \equiv a_n. \quad (40)$$

The conditional mutual information between the control action and the next population state is, in the round-indexed notation of the main text,

$$T_\pi(n) \equiv I(U_k; n_{k+1} \mid n_k = n).$$

Holding $n_k = n$ fixed, we use the shorthand $U \equiv U_k$ and $n' \equiv n_{k+1}$ below. Thus

$$I(U; n' \mid n) = \sum_{u, n'} \Pr(u, n' \mid n) \log_2 \frac{\Pr(u, n' \mid n)}{\Pr(u \mid n) \Pr(n' \mid n)}. \quad (41)$$

Using the factorization

$$\Pr(u, n' \mid n) = \Pr(u \mid n) \Pr(n' \mid u, n),$$

the factor $\Pr(u \mid n)$ cancels inside the logarithm:

$$\begin{aligned} I(U; n' \mid n) &= \sum_{u, n'} \Pr(u \mid n) \Pr(n' \mid u, n) \log_2 \frac{\Pr(n' \mid u, n)}{\Pr(n' \mid n)} \\ &= \sum_u \Pr(u \mid n) D_{\text{KL}}(\Pr(n' \mid u, n) \parallel \Pr(n' \mid n)). \end{aligned} \quad (42)$$

We therefore need the next-state distribution when the value of U is not specified. By the law of total probability,

$$\begin{aligned} \Pr(n' \mid n) &= \sum_{u \in \{0, 1\}} \Pr(u \mid n) \Pr(n' \mid u, n) \\ &= (1 - a)P(n') + aQ(n'). \end{aligned} \quad (43)$$

Define the mixture distribution

$$M(n') \equiv (1 - a)P(n') + aQ(n'). \quad (44)$$

Substituting the two possible actions into Eq. equation 42 gives

$$I(U; n' \mid n) = (1 - a)D_{\text{KL}}(P \parallel M) + aD_{\text{KL}}(Q \parallel M). \quad (45)$$

The weighted Jensen–Shannon divergence is defined by

$$\text{JS}_a(P, Q) \equiv (1 - a)D_{\text{KL}}(P \parallel M) + aD_{\text{KL}}(Q \parallel M), \quad M = (1 - a)P + aQ. \quad (46)$$

Hence, exactly,

$$I(U; n' \mid n) = \text{JS}_a(P, Q) \quad (47)$$

or, restoring the round-indexed notation,

$$T_\pi(n) \equiv I(U_k; n_{k+1} \mid n_k = n) = \text{JS}_{a_n}[Q_0(\cdot \mid n), Q_1(\cdot \mid n)]. \quad (48)$$

Up to this point, all relations are exact; no bound has yet been used.

A.1.2 PINSKER BOUND ON THE JENSEN–SHANNON DIVERGENCE

For two probability distributions A and B , define their total variation distance as

$$\text{TV}(A, B) \equiv \frac{1}{2} \sum_x |A(x) - B(x)|. \quad (49)$$

Pinsker’s inequality states that, when the Kullback–Leibler divergence is measured in bits,

$$D_{\text{KL}}(A\|B) \geq \frac{2}{\ln 2} \text{TV}(A, B)^2. \quad (50)$$

The factor $1/\ln 2$ appears because the standard Pinsker inequality is conventionally written for natural logarithms, whereas here the information is reported in bits.

Applying Eq. equation 50 separately to the two KL terms in Eq. equation 45 yields

$$T_\pi(n) \geq \frac{2}{\ln 2} [(1-a) \text{TV}(P, M)^2 + a \text{TV}(Q, M)^2]. \quad (51)$$

The two total variation distances simplify because M is a mixture of P and Q . From

$$M - P = (1-a)P + aQ - P = a(Q - P),$$

we obtain

$$\begin{aligned} \text{TV}(P, M) &= \frac{1}{2} \sum_{n'} |a(Q(n') - P(n'))| \\ &= a \text{TV}(P, Q). \end{aligned} \quad (52)$$

Similarly,

$$Q - M = Q - (1-a)P - aQ = (1-a)(Q - P),$$

and therefore

$$\text{TV}(Q, M) = (1-a) \text{TV}(P, Q). \quad (53)$$

Substituting Eqs. equation 52–equation 53 into Eq. equation 51,

$$\begin{aligned} T_\pi(n) &\geq \frac{2}{\ln 2} [(1-a)a^2 + a(1-a)^2] \text{TV}(P, Q)^2 \\ &= \frac{2}{\ln 2} a(1-a) \underbrace{[a + (1-a)]}_{=1} \text{TV}(P, Q)^2. \end{aligned} \quad (54)$$

Thus

$$T_\pi(n) \geq \frac{2a(1-a)}{\ln 2} \text{TV}(P, Q)^2. \quad (55)$$

Restoring $a = a_n$, $P = Q_0$, and $Q = Q_1$,

$$T_\pi(n) \geq \frac{2a_n(1-a_n)}{\ln 2} \text{TV}(Q_0(\cdot | n), Q_1(\cdot | n))^2. \quad (56)$$

A.1.3 FROM TOTAL VARIATION TO THE MEAN TARGET-FRACTION RESPONSE

Let $n' \in \{0, \dots, N\}$ denote the next-round number of agents supporting the controller target. Define the corresponding target fraction

$$f(n') \equiv \frac{n'}{N}. \quad (57)$$

Since $0 \leq n' \leq N$,

$$0 \leq f(n') \leq 1. \quad (58)$$

Define the intervention-induced change in the mean target fraction by

$$\begin{aligned}\chi_{h,\gamma}(n/N) &\equiv \mathbb{E}_{Q_1} \left[\frac{n'}{N} \middle| n \right] - \mathbb{E}_{Q_0} \left[\frac{n'}{N} \middle| n \right] \\ &= \mathbb{E}_Q[f] - \mathbb{E}_P[f].\end{aligned}\tag{59}$$

For completeness, we now show that any function satisfying $0 \leq f \leq 1$ obeys

$$|\mathbb{E}_Q[f] - \mathbb{E}_P[f]| \leq \text{TV}(P, Q).\tag{60}$$

Define

$$\Delta(n') \equiv Q(n') - P(n').$$

Since both P and Q are normalized,

$$\sum_{n'} \Delta(n') = 0.\tag{61}$$

Let

$$A_+ = \{n' : \Delta(n') \geq 0\}.$$

Then

$$\sum_{n' \in A_+} \Delta(n') = \text{TV}(P, Q).\tag{62}$$

Using Eq. equation 59,

$$\begin{aligned}\chi_{h,\gamma}(n/N) &= \sum_{n'} f(n') \Delta(n') \\ &= \sum_{n' \in A_+} f(n') \Delta(n') + \sum_{n' \notin A_+} f(n') \Delta(n').\end{aligned}\tag{63}$$

The second sum is non-positive because $f(n') \geq 0$ and $\Delta(n') < 0$ outside A_+ . Therefore

$$\begin{aligned}\chi_{h,\gamma}(n/N) &\leq \sum_{n' \in A_+} f(n') \Delta(n') \\ &\leq \sum_{n' \in A_+} \Delta(n') \\ &= \text{TV}(P, Q),\end{aligned}\tag{64}$$

where the second inequality uses $f(n') \leq 1$. Repeating the argument after exchanging P and Q gives the corresponding lower bound. Hence

$$|\chi_{h,\gamma}(n/N)| \leq \text{TV}(P, Q).\tag{65}$$

Squaring both sides,

$$\chi_{h,\gamma}(n/N)^2 \leq \text{TV}(P, Q)^2.\tag{66}$$

Finally, combining Eqs. equation 55 and equation 66,

$$\begin{aligned}T_\pi(n) &\geq \frac{2a_n(1-a_n)}{\ln 2} \text{TV}(Q_0(\cdot | n), Q_1(\cdot | n))^2 \\ &\geq \frac{2a_n(1-a_n)}{\ln 2} \chi_{h,\gamma}(n/N)^2.\end{aligned}\tag{67}$$

Therefore,

$$T_\pi(n) \geq \frac{2a_n(1-a_n)}{\ln 2} \chi_{h,\gamma}(n/N)^2.\tag{68}$$

For reference, the complete chain of relations is

$$\begin{aligned}
 T_\pi(n) &= I(U_k; n_{k+1} | n_k = n) \\
 &= \text{JS}_{a_n} [Q_0(\cdot | n), Q_1(\cdot | n)] \\
 &\geq \frac{2a_n(1-a_n)}{\ln 2} \text{TV}(Q_0(\cdot | n), Q_1(\cdot | n))^2 \\
 &\geq \frac{2a_n(1-a_n)}{\ln 2} \chi_{h,\gamma}(n/N)^2.
 \end{aligned} \tag{69}$$

A.2 DERIVATION OF THE THERMODYNAMIC EFFICIENCY

This appendix derives the finite-time identity and the bounded thermodynamic control efficiency stated in Section 5.5. The calculation uses only the single directional affinity h , the kinetic compliance γ , and the finite sensor-policy channel.

A.2.1 CONTROLLED TRANSITION RATIO

For one controlled microscopic opportunity,

$$K_{h,\gamma}(n+1 | n) = \gamma \frac{N-n}{N} \sigma(h), \quad K_{h,\gamma}(n | n+1) = \gamma \frac{n+1}{N} \sigma(-h).$$

Since $\sigma(h)/\sigma(-h) = e^h$,

$$\frac{K_{h,\gamma}(n+1 | n)}{K_{h,\gamma}(n | n+1)} = \frac{N-n}{n+1} e^h. \tag{70}$$

The compliance factor cancels exactly: γ changes how frequently controlled revisions occur, but not their forward/reverse directional ratio. Define the combinatorial macrostate entropy

$$S_{\text{mix}}(n) = \ln \binom{N}{n}. \tag{71}$$

Because

$$S_{\text{mix}}(n+1) - S_{\text{mix}}(n) = \ln \frac{N-n}{n+1},$$

Eq. (70) becomes

$$\ln \frac{K_{h,\gamma}(n+1 | n)}{K_{h,\gamma}(n | n+1)} = S_{\text{mix}}(n+1) - S_{\text{mix}}(n) + h. \tag{72}$$

Equivalently, the positive weight

$$w_h(n) = \binom{N}{n} e^{hn} \tag{73}$$

satisfies

$$w_h(n) K_{h,\gamma}(m | n) = w_h(m) K_{h,\gamma}(n | m) \tag{74}$$

for every supported transition. Weighted reversibility is preserved under matrix multiplication. Hence for the b -opportunity advocacy kernel $Q_1 = K_{h,\gamma}^b$,

$$\ln \frac{Q_1(m | n)}{Q_1(n | m)} = S_{\text{mix}}(m) - S_{\text{mix}}(n) + h(m-n). \tag{75}$$

For NoOp, $Q_0(m | n) = \delta_{mn}$.

A.2.2 FORWARD AND REVERSE FEEDBACK PATHS

At feedback cycle k , let $p_k(n) = P(n_k = n)$. The forward path measure on (n, y, u, m) is

$$P_F^{(k)}(n, y, u, m) = p_k(n) S(y | n) \pi(u | y) Q_u(m | n). \tag{76}$$

Its final population marginal is

$$p_{k+1}(m) = \sum_{n,y,u} P_F^{(k)}(n, y, u, m), \quad (77)$$

and the forward sensor marginal is

$$p_{Y,k}(y) = \sum_n p_k(n) S(y | n). \quad (78)$$

We use the finite-time reverse reference

$$P_R^{(k)}(m, y, u, n) = p_{k+1}(m) p_{Y,k}(y) \pi(u | y) Q_u(n | m). \quad (79)$$

The reverse reference starts from the forward final ensemble p_{k+1} and therefore does not impose stationarity. The same policy factor $\pi(u | y)$ appears in both path weights and cancels from their ratio.

Define the stochastic path irreversibility

$$\Sigma_k^* \equiv \ln \frac{P_F^{(k)}(n, y, u, m)}{P_R^{(k)}(m, y, u, n)}. \quad (80)$$

Substitution gives

$$\Sigma_k^* = \ln \frac{p_k(n)}{p_{k+1}(m)} + \ln \frac{S(y | n)}{p_{Y,k}(y)} + \ln \frac{Q_u(m | n)}{Q_u(n | m)}. \quad (81)$$

Introduce the pointwise sensing information

$$i_k(n; y) = \ln \frac{S(y | n)}{p_{Y,k}(y)}, \quad (82)$$

the stochastic system entropy

$$s_{\text{sys}}(n; p) = -\ln p(n) + S_{\text{mix}}(n), \quad (83)$$

and the controlled target-count current increment

$$j_c = u(m - n). \quad (84)$$

For $u = 0$, $m = n$ and $j_c = 0$. For $u = 1$, Eq. (75) applies. Therefore

$$\Sigma_k^* = \Delta s_{\text{sys}} + h j_c + i_k(n; y). \quad (85)$$

A.2.3 FINITE-TIME SECOND LAW AND MEAN CONTROLLED CURRENT

Averaging Eq. (85) over the forward ensemble gives

$$\Sigma_k \equiv \mathbb{E}_F[\Sigma_k^*] = D_{\text{KL}}(P_F^{(k)} \| P_R^{(k)}) \geq 0. \quad (86)$$

The three average terms are

$$\Delta S_{\text{sys},k} = S_{\text{sys}}[p_{k+1}] - S_{\text{sys}}[p_k], \quad (87)$$

$$J_{c,k} = \mathbb{E}_F[j_c], \quad (88)$$

$$I_{\text{sens},k} = I(n_k; Y_k) = \sum_{n,y} p_k(n) S(y | n) \ln \frac{S(y | n)}{p_{Y,k}(y)}, \quad (89)$$

where

$$S_{\text{sys}}[p] = -\sum_n p(n) \ln p(n) + \sum_n p(n) \ln \binom{N}{n}. \quad (90)$$

Thus

$$\Delta S_{\text{sys},k} + h J_{c,k} + I_{\text{sens},k} = \Sigma_k \geq 0. \quad (91)$$

The current can be obtained directly from the susceptibility. Conditional on $n_k = n$ and advocacy,

$$\mathbb{E}[m - n | n, U = 1] = N \chi_{h,\gamma}(n/N) = N \Lambda_{b,\gamma} \left[\sigma(h) - \frac{n}{N} \right].$$

Advocacy occurs with probability a_n , so

$$J_{c,k} = N \Lambda_{b,\gamma} \sum_n p_k(n) a_n \left[\sigma(h) - \frac{n}{N} \right]. \quad (92)$$

This is the same current entering the affinity term of Eq. (91).

A.2.4 THERMODYNAMIC CONTROL EFFICIENCY

Equation (91) contains a finite-time boundary term, $\Delta S_{\text{sys},k}$, because the population ensemble need not be stationary. To isolate the part of path irreversibility associated with the feedback operation itself, define the non-storage control expenditure

$$C_{\text{th},k} \equiv \Sigma_k - \Delta S_{\text{sys},k}. \quad (93)$$

Using Eq. (91),

$$C_{\text{th},k} = hJ_{c,k} + I_{\text{sens},k}. \quad (94)$$

The two contributions have distinct operational meanings. The first is the affinity-weighted directed population current produced by actuation; the second is the information acquired by sensing before the action is chosen.

For target-directed operation, $hJ_{c,k} \geq 0$. The fraction of the non-storage control expenditure expressed as directed current is therefore

$$\eta_{\text{th},k} \equiv \frac{hJ_{c,k}}{C_{\text{th},k}} = \frac{hJ_{c,k}}{\Sigma_k - \Delta S_{\text{sys},k}} = \frac{hJ_{c,k}}{hJ_{c,k} + I_{\text{sens},k}}. \quad (95)$$

Since $I_{\text{sens},k} \geq 0$,

$$0 \leq \eta_{\text{th},k} \leq 1 \quad (hJ_{c,k} \geq 0, C_{\text{th},k} > 0). \quad (96)$$

For a finite horizon, sum Eq. (91) over $k = 0, \dots, H-1$:

$$\sum_{k=0}^{H-1} \Delta S_{\text{sys},k} + h \sum_{k=0}^{H-1} J_{c,k} + \sum_{k=0}^{H-1} I_{\text{sens},k} = \sum_{k=0}^{H-1} \Sigma_k. \quad (97)$$

Because the system-entropy increments telescope,

$$S_{\text{sys}}[p_H] - S_{\text{sys}}[p_0] + hJ_c^{(H)} + I_{\text{sens}}^{(H)} = \Sigma^{(H)}, \quad (98)$$

where

$$J_c^{(H)} \equiv \sum_{k=0}^{H-1} J_{c,k}, \quad I_{\text{sens}}^{(H)} \equiv \sum_{k=0}^{H-1} I_{\text{sens},k}, \quad \Sigma^{(H)} \equiv \sum_{k=0}^{H-1} \Sigma_k. \quad (99)$$

The accumulated non-storage expenditure and corresponding efficiency are therefore

$$C_{\text{th}}^{(H)} = hJ_c^{(H)} + I_{\text{sens}}^{(H)}, \quad \eta_{\text{th}}^{(H)} = \frac{hJ_c^{(H)}}{C_{\text{th}}^{(H)}}. \quad (100)$$

In general,

$$\eta_{\text{th}}^{(H)} \neq \frac{1}{H} \sum_{k=0}^{H-1} \eta_{\text{th},k}, \quad (101)$$

because thermodynamic efficiency is a nonlinear ratio and the cycle-wise expenditures need not be equal.

At stationarity the boundary term vanishes and

$$\eta_{\text{th}} = \frac{hJ_c}{\Sigma}. \quad (102)$$

The finite-time expression is therefore not obtained by silently assuming $\Delta S_{\text{sys}} = 0$; rather, it removes that transient storage term explicitly before forming the ratio.

The efficiency makes the sensing–actuation trade-off transparent. Holding the sensor channel fixed, increasing actuation can raise hJ_c and drive η_{th} toward one. Conversely, increasing q_c can increase I_{sens} without a proportional increase in J_c , thereby reducing η_{th} . Kinetic compliance enters through J_c : a controller can have a large directional affinity but low thermodynamic efficiency when agents rarely comply and the resulting current remains small relative to the sensing information acquired.

B RELATIONAL-REASONING PROMPT

The population is driven by one versioned prompt family, `relational_public_ballot` version 1. Each call is compiled into two chat messages. The system message establishes the agent’s role, privacy boundary, task, call-specific option mapping, and currently known facts. The user message supplies the current vote and visible social source, then states the response contract. The model therefore receives only the state available to the focal agent at that microscopic update; it is never shown the global vote vector, other agents’ private facts, the ground-truth answer, or the controller’s sensor sample.

The display below reproduces one complete prompt from the audit of the minimal Study 04 inspection episode ($N = 24$, $q_c = 12$, $b = 12$). It is the first focal update of population round 1 for Agent 17. Typography is annotation only: **orange** marks values instantiated anew for the call, **green** marks the focal agent’s private state, and **blue** marks the visible social input. None of these colours or the “SYSTEM MESSAGE”/“USER MESSAGE” labels are sent to the model.

Audited focal-update prompt

SYSTEM MESSAGE

You are **Agent 17**, one participant in a group reasoning problem.

Your goal is to identify the correct answer.

Participant identities persist across interactions.

Different participants know different facts about this problem. No single participant is guaranteed to know everything needed to answer, and some participants may have objectives that differ from yours. You do not know which participants, if any, have different objectives.

Make your own decision, using:

- the facts you currently know;
- your current position;
- the public positions of the participants shown below.

You see only the participants shown in this interaction, and of each one you see their vote and the fact they chose to share, if any.

Other participants will see the same of you: your vote and any fact you choose to share, and nothing else.

Your reason is your own record: it is not shown to anyone.

QUESTION

Where is **Davi** relative to **Hira**?

The available answers are:

- **A) WEST**
- **B) EAST**
- **C) SOUTHWEST**

Exactly one of these answers is correct. Vote by its letter.

YOUR CURRENT KNOWLEDGE

You do not currently know any facts about this problem. You will have to rely on what other participants tell you.

USER MESSAGE

YOUR CURRENT POSITION

Vote: **SOUTHWEST**

CURRENT SOCIAL INFORMATION

Agent **25**

Vote: **B**

DECISION

Work out which option the facts available to you support, and vote for it.

Your reason should briefly explain your choice, for your own record.

Sharing a fact is the only way to pass information to other participants. You may share exactly one of the facts you currently know by giving its identifier in `shared_fact_id`, so that the participants who see your position can use it too. Use “none” if you prefer to share nothing.

You may share only a fact listed under YOUR CURRENT KNOWLEDGE, by its exact identifier. Do not invent facts, identifiers, or relationships that were not given to you.

Keep your reason to at most three sentences.

Return only valid JSON:

```
{
  "vote": "<A | B | C>",
```

```

"reason": "<brief private reason>",
"shared_fact_id": "<none>"
}

```

What changes between calls. The identity, query, available semantic answers, private facts, current vote, visible sources, and admissible fact identifiers are bound from the live game state. The semantic answers are independently permuted onto A/B/C for every call. Thus the *social vote B* above means *EAST* only in this instance; the persistent population state stores *EAST*, not the letter *B*. This prevents a stable display position from becoming an unintended convention.

Ordinary and controlled social inputs. An uncontrolled update shows the vote and, when supplied, one valid fact from an ordinary peer. In the audited example the visible source is the controller: the controller’s semantic target *EAST* is rendered as *Vote: B* under this call’s permutation. The prompt deliberately does not identify that source as a controller or authority. Under the *recommendation_only* mode it also supplies no fact, so the intervention can change the public vote state but cannot directly enlarge the focal agent’s knowledge set.

Initialization variant. Before social interaction, each of the 24 agents receives the same role, task, privacy, fact-sharing, and JSON-contract blocks, but no current-position or social-information block. Instead, the prompt states that no other participant has yet declared a position and instructs the agent to decide from its local facts alone. The resulting 24 ballots define the initial population state.

Response and information boundary. The accepted response contains a display-letter vote, a private reason of at most three sentences, and either one fact identifier already in the agent’s knowledge set or *none*. The display letter is immediately mapped back to its semantic relation. The reason is retained for audit but never shown to another agent. Only the semantic vote and the explicitly selected fact may enter a later prompt, which makes the public decision channel and the epistemic transmission channel separately observable.

Audit check. The inspection episode recorded 72 distinct compiled prompt instances: 24 local-initialization prompts and 48 focal-update prompts (24 in each of two population rounds). All 72 responses passed validation without a retry or provider error. The prompt definition hash was constant across the episode, whereas every instance hash was unique, confirming that one fixed template was bound to 72 call-specific views.

B.1 EPISTEMIC CONDITIONS: RECEIVER DISPOSITION AND CONTROLLER EVIDENCE

The population game and the controller each expose one binary epistemic knob, and the two are set independently. The first, *receiver_epistemic_disposition*, decides what every agent is told about the social channel it listens to. The second, *controller_evidence_strategy*, decides which single fact the controller is allowed to cite when it speaks. They live at *game.options* and *control.options* respectively. Crossing them gives the four conditions used throughout Study 08, recorded in each round record as *derived_epistemic_condition*:

Condition	Receiver disposition	Controller evidence
<i>naive_neutral</i>	<i>naive</i>	<i>neutral</i>
<i>naive_strategic</i>	<i>naive</i>	<i>strategic</i>
<i>vigilant_neutral</i>	<i>vigilant</i>	<i>neutral</i>
<i>vigilant_strategic</i>	<i>vigilant</i>	<i>strategic</i>

The two axes act at different places in the loop. The receiver axis is a *text* intervention: it rewrites one fixed system block, identically for every agent and every call of an episode. The controller axis is a *content* intervention: it changes no template text at all, only which proposition of the frozen task the controller puts into its social slot. An agent therefore cannot read the controller’s strategy off the wording of its prompt, and the two manipulations remain separately attributable.

Receiver epistemic disposition. This axis substitutes the `social_environment` block, the second system block of the manifest. Both variants are bound as *fixed*: they are part of the episode’s prompt definition rather than of any per-call view, so within one cell every one of the episode’s prompts carries the same wording. *naive* states only that identities persist, and says nothing about whether a visible participant is worth believing. *vigilant* adds an explicit instruction to discount a bare recommendation and to check the evidence actually in hand. The two blocks are reproduced below exactly as sent.

social_environment — naive

Participant identities persist across interactions.

social_environment — vigilant

Participant identities persist across interactions.

Some participants may have objectives that differ from yours. Do not accept a recommendation merely because another participant gives it. Evaluate whether the explicit evidence you actually have supports the conclusion, and remain aware that a source may present only part of the relevant information.

Nothing else in the prompt moves with this axis. In the audited grid the block is estimated at 6 tokens under *naive* and 55 under *vigilant*, carrying the whole per-prompt difference of 160 versus 209 estimated tokens. The block manifest — names, order, roles, and bindings — is byte-identical across all four cells, because it records prompt *structure*; the condition is pinned by the `receiver_epistemic_disposition` field of each cell’s resolved configuration and is visible in the rendered text itself.

The audited Study 04 prompt displayed earlier in this appendix predates this factorization. It shows the older combined block, which asserts both distributed information and possible objective divergence in one paragraph. The factorized axis separates those two claims, and Study 08 uses only the *naive/vigilant* pair above.

Controller evidence strategy. Evidence reaches the population only in the controller’s message mode `recommendation_plus_fact`, in which the controller occupies its social slot with both a vote and exactly one fact of the frozen task. That fact is chosen once, before the episode runs, by a deterministic rule that never consults a model:

- *neutral* takes the first fact in the task’s frozen fact order. The choice ignores which relation the controller advocates, so the cited evidence is fixed by the task alone.
- *strategic* takes the admissible fact whose relation points most nearly in the advocated direction. Each of the eight compass relations has a symbolic displacement vector; the controller scores every fact by the cosine between its relation vector and the advocated relation’s vector, discards every fact scoring zero or below, and cites the survivor with the highest score, breaking ties toward the earlier fact. If no fact has positive alignment the configuration is rejected at launch rather than silently advocating without evidence.

strategic is thus the selective-evidence arm: the controller still never fabricates anything, since only real facts of the task are admissible, but it is free to cite the true statement that best flatters its target and to stay silent about the rest. *neutral* removes that freedom while holding the recommendation, the budget, and the number of exposures fixed, so the difference between the arms isolates the value of evidence *selection*.

Worked example. Study 08’s task `task_0001` asks where Zora is relative to Feni, with frozen fact order (`f1`, `f2`, `f3`): `f1` “Zora is south of Maro” (SOUTH), `f2` “Maro is west of Feni” (WEST), and the distractor `f3` “Garó is north of Belo” (NORTH), which mentions neither query entity. The correct relation is SOUTHWEST. The rules then resolve as follows.

Advocated relation	neutral	strategic
SOUTHWEST (correct)	f1 (SOUTH)	f1 (SOUTH)
NORTHWEST	f1 (SOUTH)	f2 (WEST)
NORTH	f1 (SOUTH)	f3 (NORTH)

The last row is the informative one. Advocating NORTH, the neutral controller must cite f1, whose southward relation openly contradicts its own recommendation. The strategic controller instead cites f3: a true statement, pointing north, about two entities irrelevant to the question. Under *naive* the agent is given no reason to notice the substitution; under *vigilant* it is told in so many words to check whether the evidence it holds supports the conclusion.

Audited grid. The four conditions were run once end to end against a live model. Each cell is one episode of $N = 6$ agents over 4 population rounds with a single visible social slot ($K = 1$), sensor sample $q_c = 3$, exact budget $b = 3$, target *correct*, *recommendation_plus_fact*, on `gwdg/openai-gpt-oss-120b` at temperature 0 with seed 20260828. Each cell issued 30 calls — 6 local-initialization ballots and $6 \times 4 = 24$ focal updates — for 120 in total, matching the recorded request count exactly, so no response required a retry.

In this particular task the two evidence arms coincide: the advocated relation is the correct one, SOUTHWEST, and the tie between f1 and f2 — both at cosine $1/\sqrt{2}$ — resolves to the earlier fact, which is also the neutral rule’s choice. All four cells therefore exposed f1, and the evidence axis is degenerate here by construction. The grid is a prompt- and plumbing-level audit of the four conditions, not an estimate of their effect; the axis separates only on tasks and targets where alignment discriminates, as the two lower rows of the worked example show.

C CONTROLLED SUSCEPTIBILITY STUDY

C.1 STATE-MATCHED RESPONSE TO ADVOCACY

We isolate the one-round behavioral response of the population to controller advocacy while holding the epistemic population fixed. The experiment uses $N = 24$ agents, $q = 1$ social interaction, reasoning depth $L = 2$, support redundancy $r = 6$, recommendation-only control, and an incorrect controller target Z . We fix the sensing and actuation budgets to $q_c = b = 18$ and explicitly initialize three target-support levels,

$$x_0 \in \{0.25, 0.375, 0.50\}, \quad n_Z(0) \in \{6, 9, 12\}.$$

Within each task, agent identities and initial knowledge sets $K_i(0)$ are identical across all three conditions, giving $\kappa_0 = 0.25$ and $\phi_0 = 0$. Target supporters are selected deterministically and balanced across the populated knowledge strata, with nested sets $Z_6 \subset Z_9 \subset Z_{12}$; all remaining agents are initialized at the correct relation. Thus only the initial social support for Z changes across conditions.

For each x_0 , we compare two deterministic one-round interventions: *ADVOCATE*, in which the controller advocates for Z at exactly $b = 18$ controlled update positions, and *NOOP*, in which it does not intervene. The two arms use matched episode seeds. We evaluate two tasks with ten matched repetitions per task, yielding 120 one-round episodes. The sensor remains active for logging, but it does not determine the forced action.

The primary quantity is the signed intervention response

$$\chi(x_0) = \mathbb{E}[x_1 - x_0 \mid \text{ADV}, x_0] - \mathbb{E}[x_1 - x_0 \mid \text{NOOP}, x_0], \quad (103)$$

with 95% confidence intervals from 1000 bootstrap resamples of matched episode pairs. Since the action is deterministic within each arm, this experiment estimates a response contrast rather than an action-to-population mutual information.

The measured responses are $\chi(0.25) = 0.163$ (95% CI $[0.119, 0.202]$), $\chi(0.375) = 0.100$ ($[0.060, 0.144]$), and $\chi(0.50) = 0.104$ ($[0.060, 0.148]$). The controller therefore produces a positive causal shift toward its target at all three matched states. However, the largest response occurs at the lowest tested target fraction, and increasing the initial target support from $x_0 = 0.375$ to 0.50




Figure 5: One-round state-matched response to advocacy. Points show the empirical ADVOCATE-minus-NOOP response with 95% matched-pair bootstrap intervals. The dashed curve is the exact $q = 1$ reference at $N = 24$ and $b = 18$, $\chi_{q=1}(x_0) = (1 - x_0)[1 - (1 - 1/N)^b]$. Advocacy produces a positive response at every imposed initial target fraction, while the response does not increase with pre-existing target support over the tested range.

produces no additional gain. The exact $q = 1$ reference predicts the same qualitative ordering—largest leverage at low target support—but substantially larger responses (0.401, 0.334, and 0.268, respectively). Thus, in this controlled one-round setting, LLM advocacy is effective but markedly weaker than the idealized copying reference, and there is no evidence that additional pre-existing social support amplifies the intervention response.

C.2 EMPIRICAL CALIBRATION OF DIRECTIONAL AFFINITY AND COMPLIANCE

The single-affinity theory describes a controlled interaction using a directional bias h and a kinetic compliance γ . To connect these quantities to the LLM dynamics, we use the microscopic transitions observed in the state-matched intervention experiment. We coarse-grain each focal agent’s vote into two states: supporting the controller target Z , or not supporting Z . Among controlled updates, define

$$p_+ = P(\text{non-}Z \rightarrow Z \mid \text{controlled}), \quad p_- = P(Z \rightarrow \text{non-}Z \mid \text{controlled}).$$

Under the finite-compliance kernel,

$$p_+ = \gamma \sigma(h), \quad p_- = \gamma \sigma(-h).$$

Therefore the two parameters are directly identifiable:

$$\gamma_{\text{eff}} = p_+ + p_-, \quad h_{\text{eff}} = \ln \frac{p_+}{p_-}. \quad (104)$$

The sum measures how readily a controlled opportunity produces any revision, whereas the ratio measures the direction favored conditional on revision.

Pooling all controlled updates gives

$$h_{\text{eff}} = 3.83, \quad \gamma_{\text{eff}} = 0.372,$$

Table 1: Microscopic transitions observed at controlled update positions.

Transition	Eligible updates	Observed transitions	Probability
non- $Z \rightarrow Z$	572	208	$p_+ = 0.364$
$Z \rightarrow$ non- Z	508	4	$p_- = 0.0079$

Table 2: Empirical quantities connecting microscopic LLM dynamics to the single-affinity finite-compliance description.

Quantity	Estimate	Interpretation
h_{eff}	3.83	directional affinity
95% CI for h_{eff}	[3.06, 4.88]	bootstrap uncertainty
$e^{h_{\text{eff}}}$	46.2	forward/reverse transition ratio
γ_{eff}	0.372	probability scale for actual revision

with forward-to-reverse transition ratio $e^{h_{\text{eff}}} \simeq 46.2$. Thus motion toward the controller target is strongly favored among revisions, while most controlled opportunities still leave the focal vote unchanged. A stratified episode-level bootstrap gives the approximate 95% interval

$$h_{\text{eff}} \in [3.06, 4.88].$$

No further decomposition of h_{eff} is required by the finite-time thermodynamic theory. The same measured affinity enters the susceptibility, the controlled current, and the path identity $\Delta S_{\text{sys}} + hJ_c + I(n_k; Y_k) = \Sigma$.

D RESOURCE-CONSTRAINED CONTROL DESIGN

The theory in Section 5 separates three questions that are often conflated in controller design: how strongly the population can be moved, how coherently controller influence is converted into that motion, and how efficiently the finite-time sensing-actuation budget is converted into directed current. This appendix turns susceptibility, the two bounded efficiencies, and an explicit operational interaction cost into a practical design prescription.

D.1 MINIMUM-COST CONTROL AT PRESCRIBED RESPONSE

Let

$$\varphi = (q_c, b, \theta, \beta, \dots)$$

denote the tunable feedback-controller configuration, and let $\rho(n)$ be the population-state occupancy over the evaluation window. The microscopic channel parameters h and γ are treated separately: unless the implementation exposes a mechanism that changes them directly, they are calibrated properties of the population response rather than free design knobs.

Since the state-local susceptibility is defined conditionally on advocacy, the expected controller-induced response per round is

$$R(\varphi; h, \gamma) = \sum_{n=0}^N \rho(n) a_n(\varphi) \chi_{h,\gamma} \left(\frac{n}{N} \right). \quad (105)$$

A direct operational interaction cost is

$$C_{\text{op}}(\varphi) = q_c + b \sum_{n=0}^N \rho(n) a_n(\varphi), \quad (106)$$

where the first term counts the agents sensed every round and the second counts the expected number of controlled microscopic opportunities.

The most direct controller-design problem is therefore not to maximize an efficiency in isolation, but to achieve a prescribed behavioral response at minimum operational cost,

$$\varphi^* = \arg \min_{\varphi} C_{\text{op}}(\varphi) \quad \text{subject to} \quad R(\varphi; h, \gamma) \geq R_{\text{req}}. \quad (107)$$

Equivalently, for a fixed interaction budget one may maximize the achievable response,

$$\varphi^* = \arg \max_{\varphi} R(\varphi; h, \gamma) \quad \text{subject to} \quad C_{\text{op}}(\varphi) \leq C_{\text{max}}. \quad (108)$$

These formulations make the practical question explicit: once the required response has been fixed, how much sensing and actuation can be removed before performance falls below the desired level?

For the exact finite-compliance $q = 1$ reference,

$$\chi_{h,\gamma}(x) = [\sigma(h) - x] \Lambda_{b,\gamma}, \quad \Lambda_{b,\gamma} = 1 - \left(1 - \frac{\gamma}{N}\right)^b. \quad (109)$$

The marginal gain from one additional controlled opportunity is

$$\Lambda_{b+1,\gamma} - \Lambda_{b,\gamma} = \frac{\gamma}{N} \left(1 - \frac{\gamma}{N}\right)^b, \quad (110)$$

which decreases monotonically with b . Thus stronger actuation continues to increase the mean response toward the set point, but each additional controlled opportunity contributes less than the previous one. The same logic applies to sensing: increasing q_c is useful only insofar as the improved state estimate changes the policy in a way that improves the required response.

D.2 HOW THE TWO EFFICIENCIES GUIDE FEASIBLE DESIGNS

The operational problem above specifies how to purchase less control. The two bounded efficiencies from Section 5 then determine whether a feasible controller uses its influence and thermodynamic budget well.

The information–response efficiency is

$$\eta_{\text{IR}}(n) = \frac{2a_n(1 - a_n) \chi_{h,\gamma}(n/N)^2}{(\ln 2) T_{\pi}(n)} \leq 1, \quad (111)$$

with $T_{\pi}(n) = I(U_k; n_{k+1} | n_k = n)$. It measures whether the action-dependent change in the full next-state distribution is expressed as a coherent shift in mean target fraction. Among settings with comparable cost and response, one should prefer regimes in which η_{IR} remains high over the states actually visited.

The thermodynamic efficiency is

$$\eta_{\text{th},k} = \frac{hJ_{c,k}}{\Sigma_k - \Delta S_{\text{sys},k}} = \frac{hJ_{c,k}}{hJ_{c,k} + I(n_k; Y_k)} \leq 1. \quad (112)$$

Here $\Delta S_{\text{sys},k}$ is the transient storage term and $\Sigma_k - \Delta S_{\text{sys},k}$ is the non-storage finite-time control expenditure. The efficiency therefore reports the fraction of the generalized sensing–actuation expenditure expressed as affinity-weighted directed current. Its derivation from the forward/reverse path ratio is given in Appendix A.2.

Two practical rules follow directly. First, more sensing is not automatically better: if increasing q_c raises $I(n_k; Y_k)$ without a proportional increase in $hJ_{c,k}$, it lowers η_{th} . Second, more actuation is useful only while it continues to raise the required response; the exact susceptibility already shows diminishing marginal gains through $\Lambda_{b+1,\gamma} - \Lambda_{b,\gamma}$.

Taken together, the design prescription is

- (i) specify the minimum acceptable response or current,
- (ii) minimize operational cost subject to that requirement,
- (iii) use η_{IR} to reject incoherent action channels,
- (iv) among feasible controllers, prefer larger η_{th} ,
- (v) retain h and γ as separate directional and kinetic diagnostics.

Thus susceptibility determines where control has leverage, the operational cost determines how much control must be purchased, η_{IR} determines whether controller influence becomes coherent mean motion, and η_{th} determines whether sensing and actuation are converted efficiently into directed current.

D.3 FINITE-TIME AND NONSTATIONARY INTERPRETATION

No stationarity assumption is required in the construction above. Let

$$\rho_k(n) = \Pr(n_k = n) \quad (114)$$

denote the transient population distribution at round k . Over a finite horizon of H rounds, the occupancy entering Eqs. (105)–(106) can be interpreted as the finite-time visitation distribution

$$\bar{\rho}_H(n) = \frac{1}{H} \sum_{k=0}^{H-1} \rho_k(n). \quad (115)$$

The visitation distribution $\bar{\rho}_H$ can be used directly for observables that are linear in the occupancy, such as the expected controlled response and current. For example,

$$\frac{1}{H} \sum_{k=0}^{H-1} J_{c,k} = N \sum_n \bar{\rho}_H(n) a_n \chi(n). \quad (116)$$

Sensing information is different because mutual information is nonlinear in the state distribution. In general,

$$I_{\text{sens}}[\bar{\rho}_H] \neq \frac{1}{H} \sum_{k=0}^{H-1} I_{\text{sens}}[\rho_k]. \quad (117)$$

Finite-horizon thermodynamic accounting therefore evaluates $I_{\text{sens},k}$ from each transient distribution ρ_k and accumulates the resulting cycle-wise contributions before forming $\eta_{\text{th}}^{(H)}$. This preserves the finite-time interpretation without imposing stationarity or replacing the transient sequence by a single pooled ensemble.

Hence the operational design problem averages over the transient states actually visited during the experiment rather than over an equilibrium distribution. Likewise, Eq. (91) retains the finite-time boundary contribution $\Delta S_{\text{sys},k}$, which accounts for the fact that the population ensemble changes during the feedback cycle. A stationary distribution is therefore only a special long-time limit of the present formulation, not an assumption required by it.